

Simplify Machine Learning With Flink TableAPI

公司：阿里巴巴

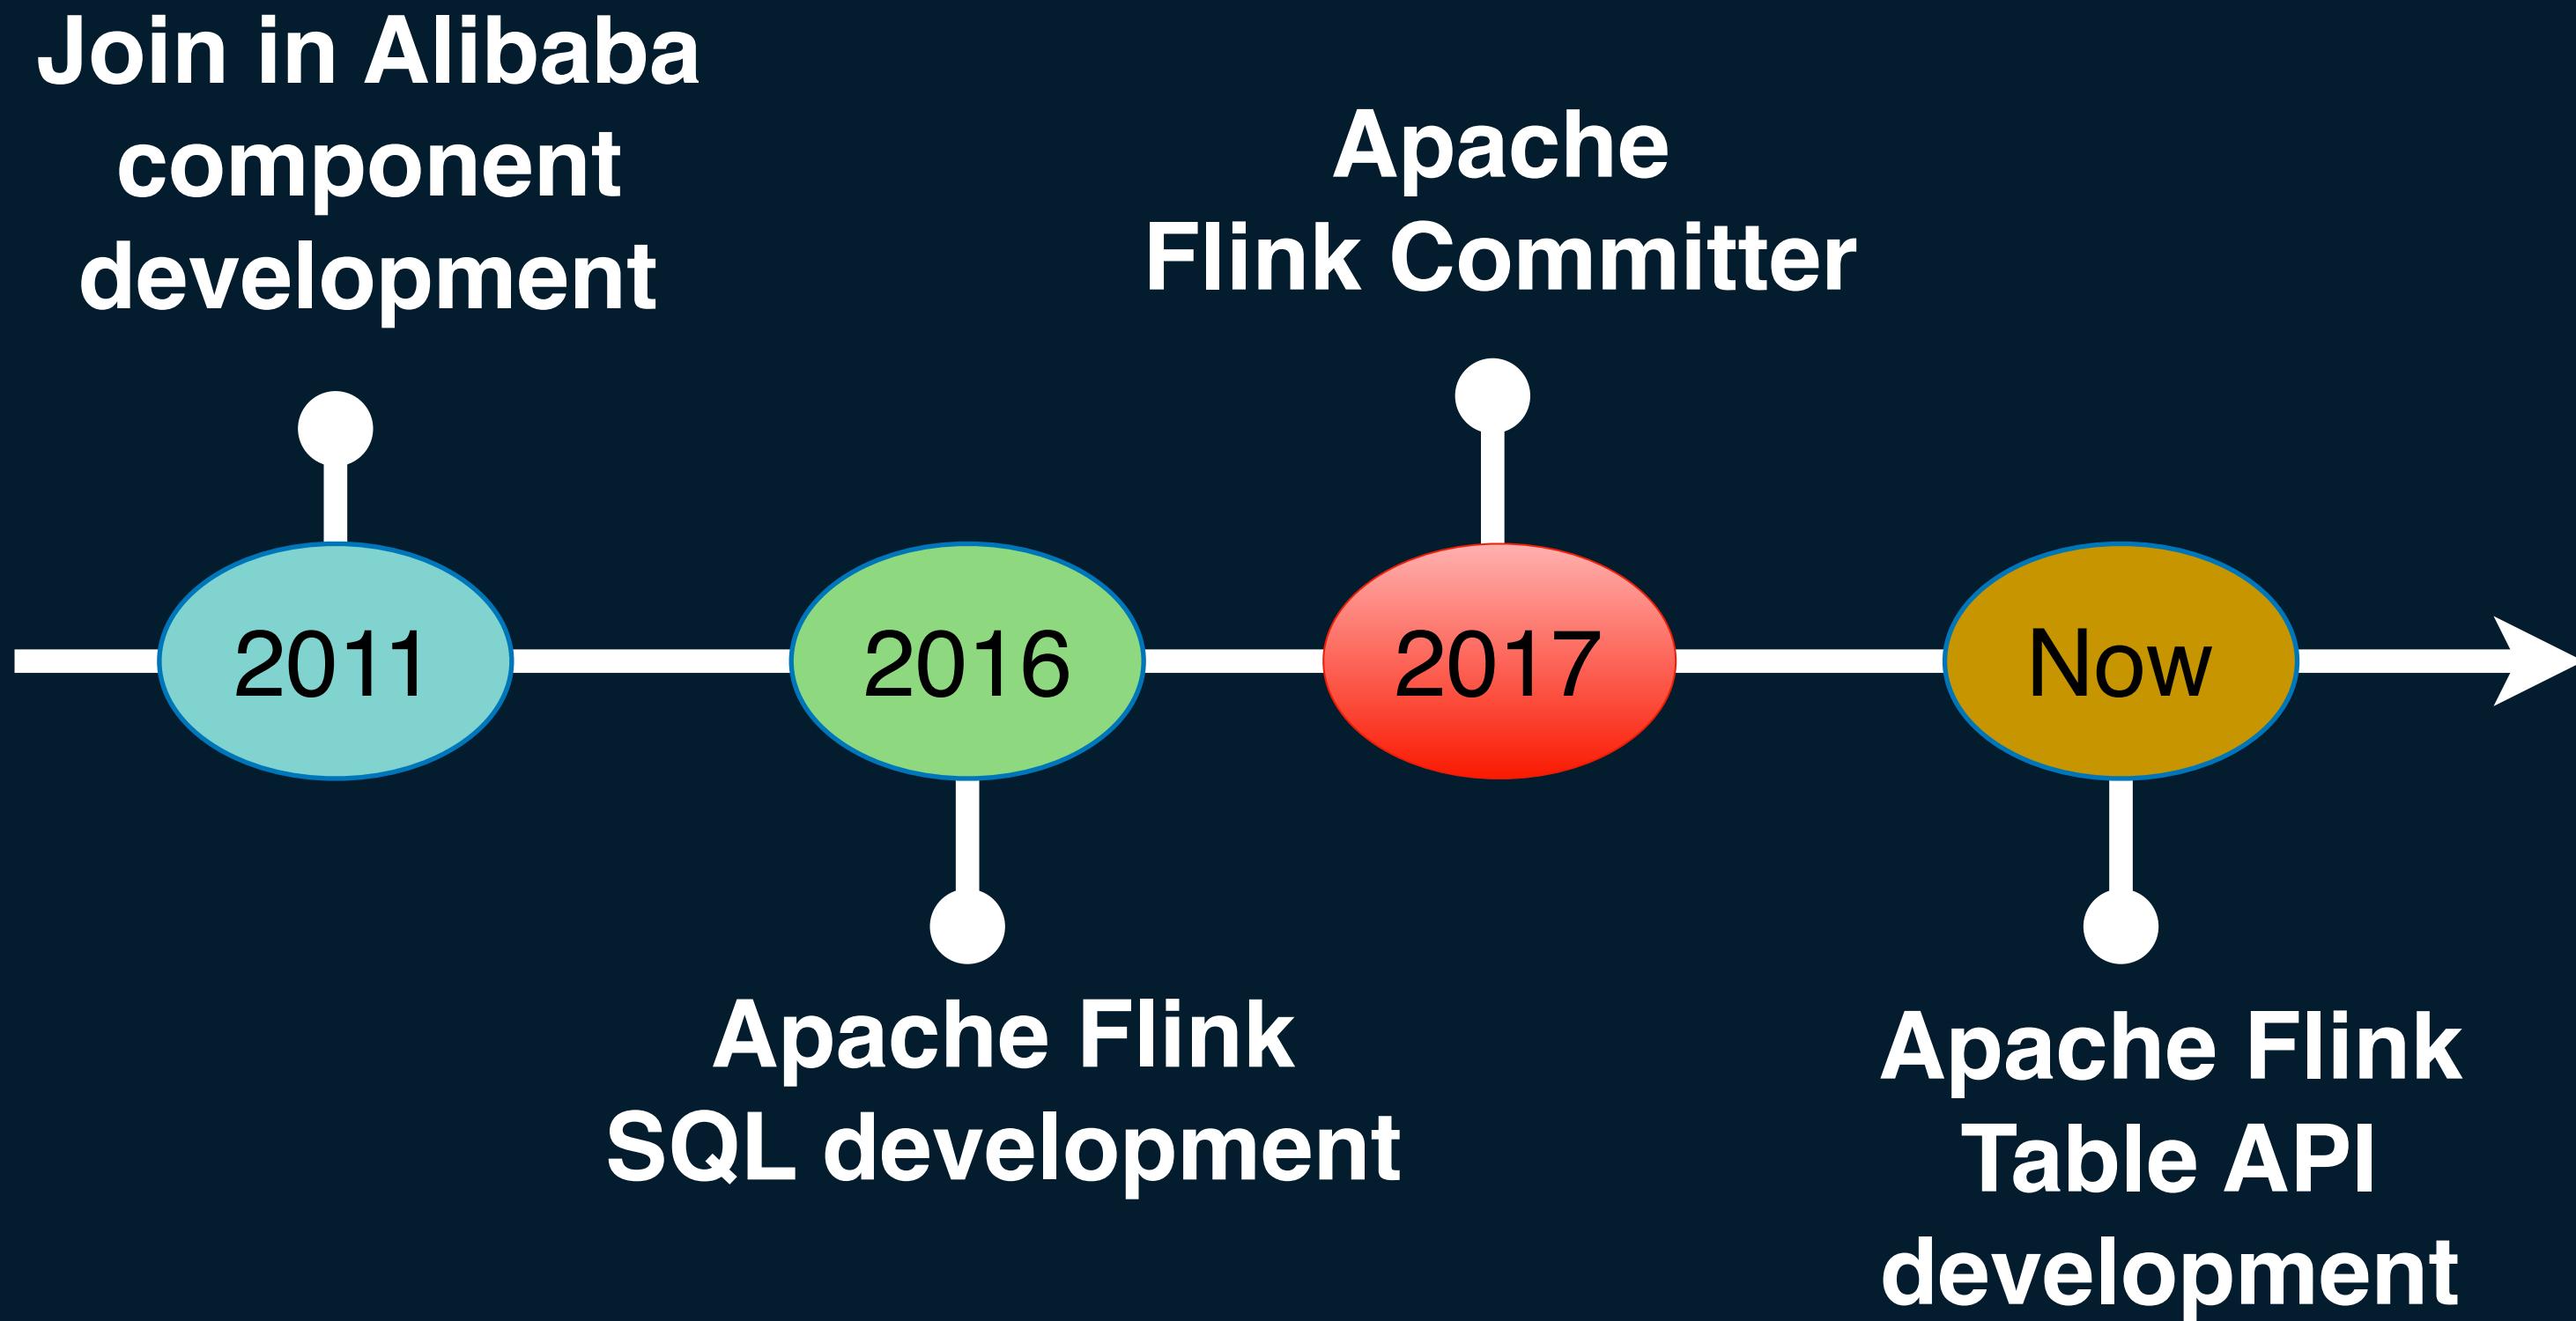
演讲者：孙金城（高级技术专家）

秦江杰（高级技术专家）





Jincheng Sun(JinZhu)
孙金城(金竹)





Jiangjie (Becket) Qin
秦江杰

Alibaba Realtime Computing Platform, 2018.5 -

2018年5月加入阿里巴巴实时计算平台

Worked at LinkedIn and IBM

曾就职于 LinkedIn 和 IBM

M.S Information Networking, CMU, 2014

2014年硕士毕业于卡耐基梅陇大学

Apache Kafka PMC Member

Apache Kafka 项目管理委员会成员

AGENDA

提要



Apache Flink

Brief introduction to Flink Table API

Table API 简介

The API requirements from Machine Learning Algorithms

机器学习算法对 API 的核心需求

The enhancement to Flink Table API

Flink Table API 的扩展

Algorithm examples based on Flink Table API

基于 Flink Table API 的算法实现

AGENDA

提要



Apache Flink

Brief introduction to Flink Table API

Table API 简介

The API requirements from Machine Learning Algorithms

机器学习算法对 API 的核心需求

The enhancement to Flink Table API

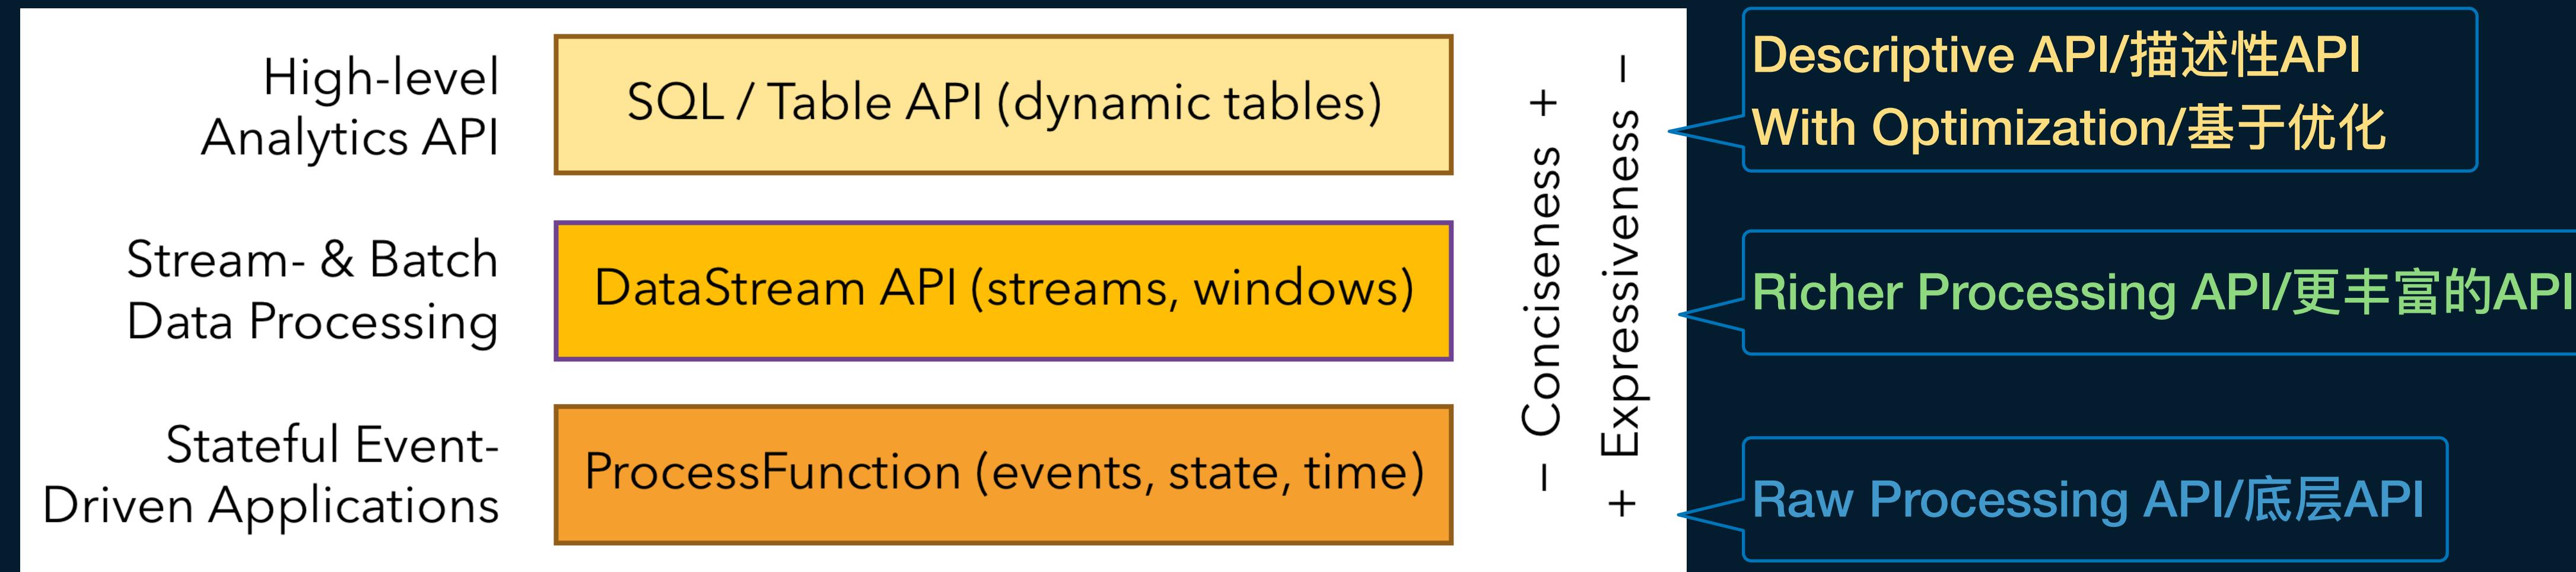
Flink Table API 的扩展

Algorithm examples based on Flink Table API

基于 Flink Table API 的算法实现

Table API is high-level analytics API

Apache Flink Layered APIs



E.g.: Count the numbers of people by Region

示例：按地区统计人口数量

table
.groupBy('region')
.select('region, COUNT(1))

```
def main(args: Array[String]): Unit = {
    val env = StreamExecutionEnvironment.getExecutionEnvironment
    env.setParallelism(1)

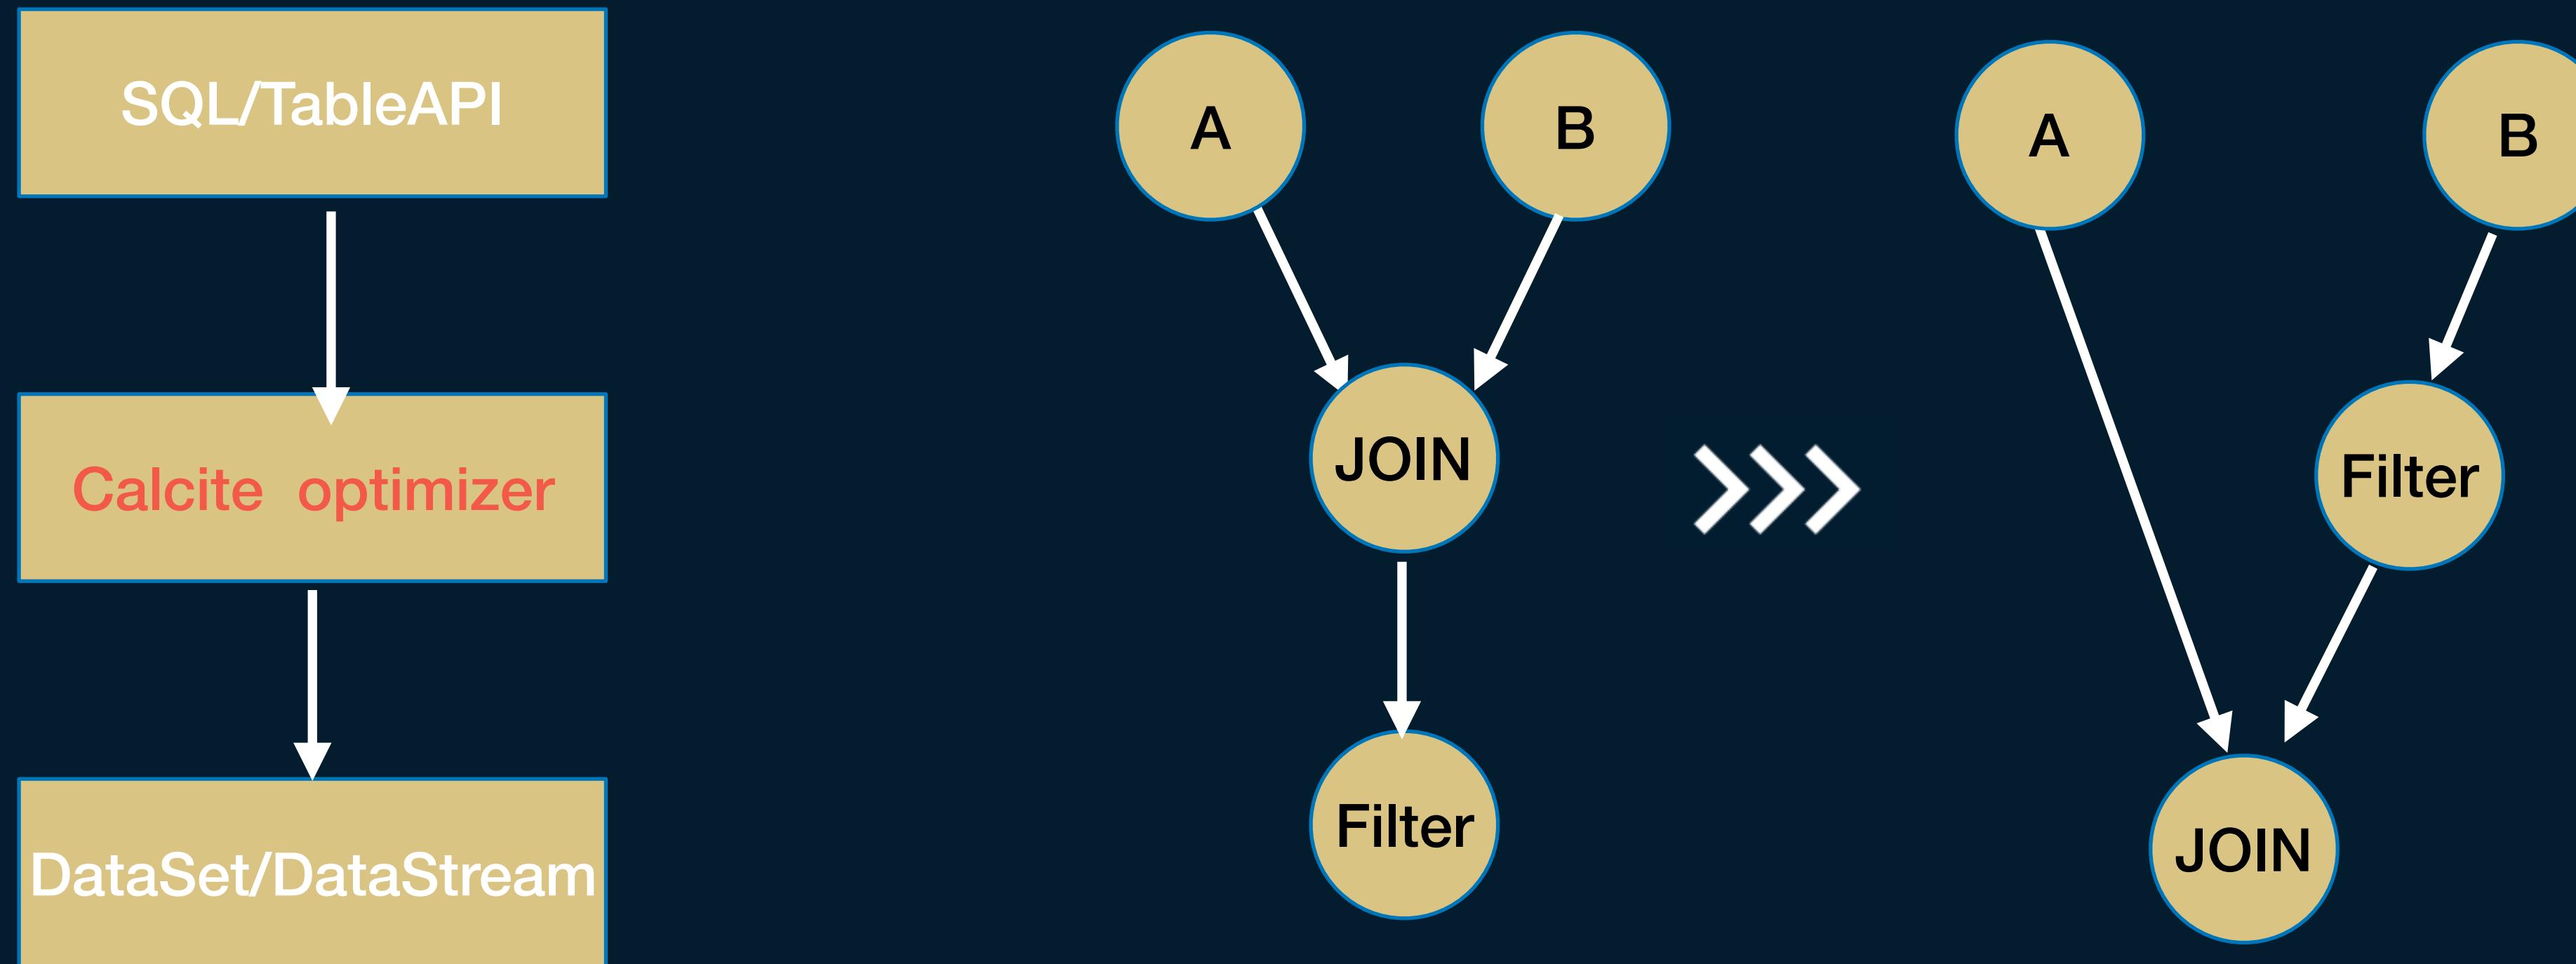
    val dataStream = env.addSource(new WordCountSourceFunction)
    val tokenizerWord = dataStream.flatMap(new TokenizerFlatMap)
    val result = tokenizerWord
        .map{ _, 1}
        .keyBy( fields = 0)
        .sum( position = 1)

    result.addSink(new PrintSinkFunction[(String, Int)])
    env.execute()
}

class TokenizerFlatMap extends FlatMapFunction[String, String] {
    override def flatMap(
        t: String,
        out: Collector[String]): Unit = {
        val stringTokenizer = new StringTokenizer(t, ",!`.;", returnDelims = false)
        while(stringTokenizer.hasMoreElements()) {
            val word = stringTokenizer.nextToken().trim.toLowerCase
            out.collect(word)
        }
    }
}
```

10+

TableAPI is optimized



Get better execution performance
获取更好执行性能

TableAPI unifies batch and stream

Component Stack 组件栈

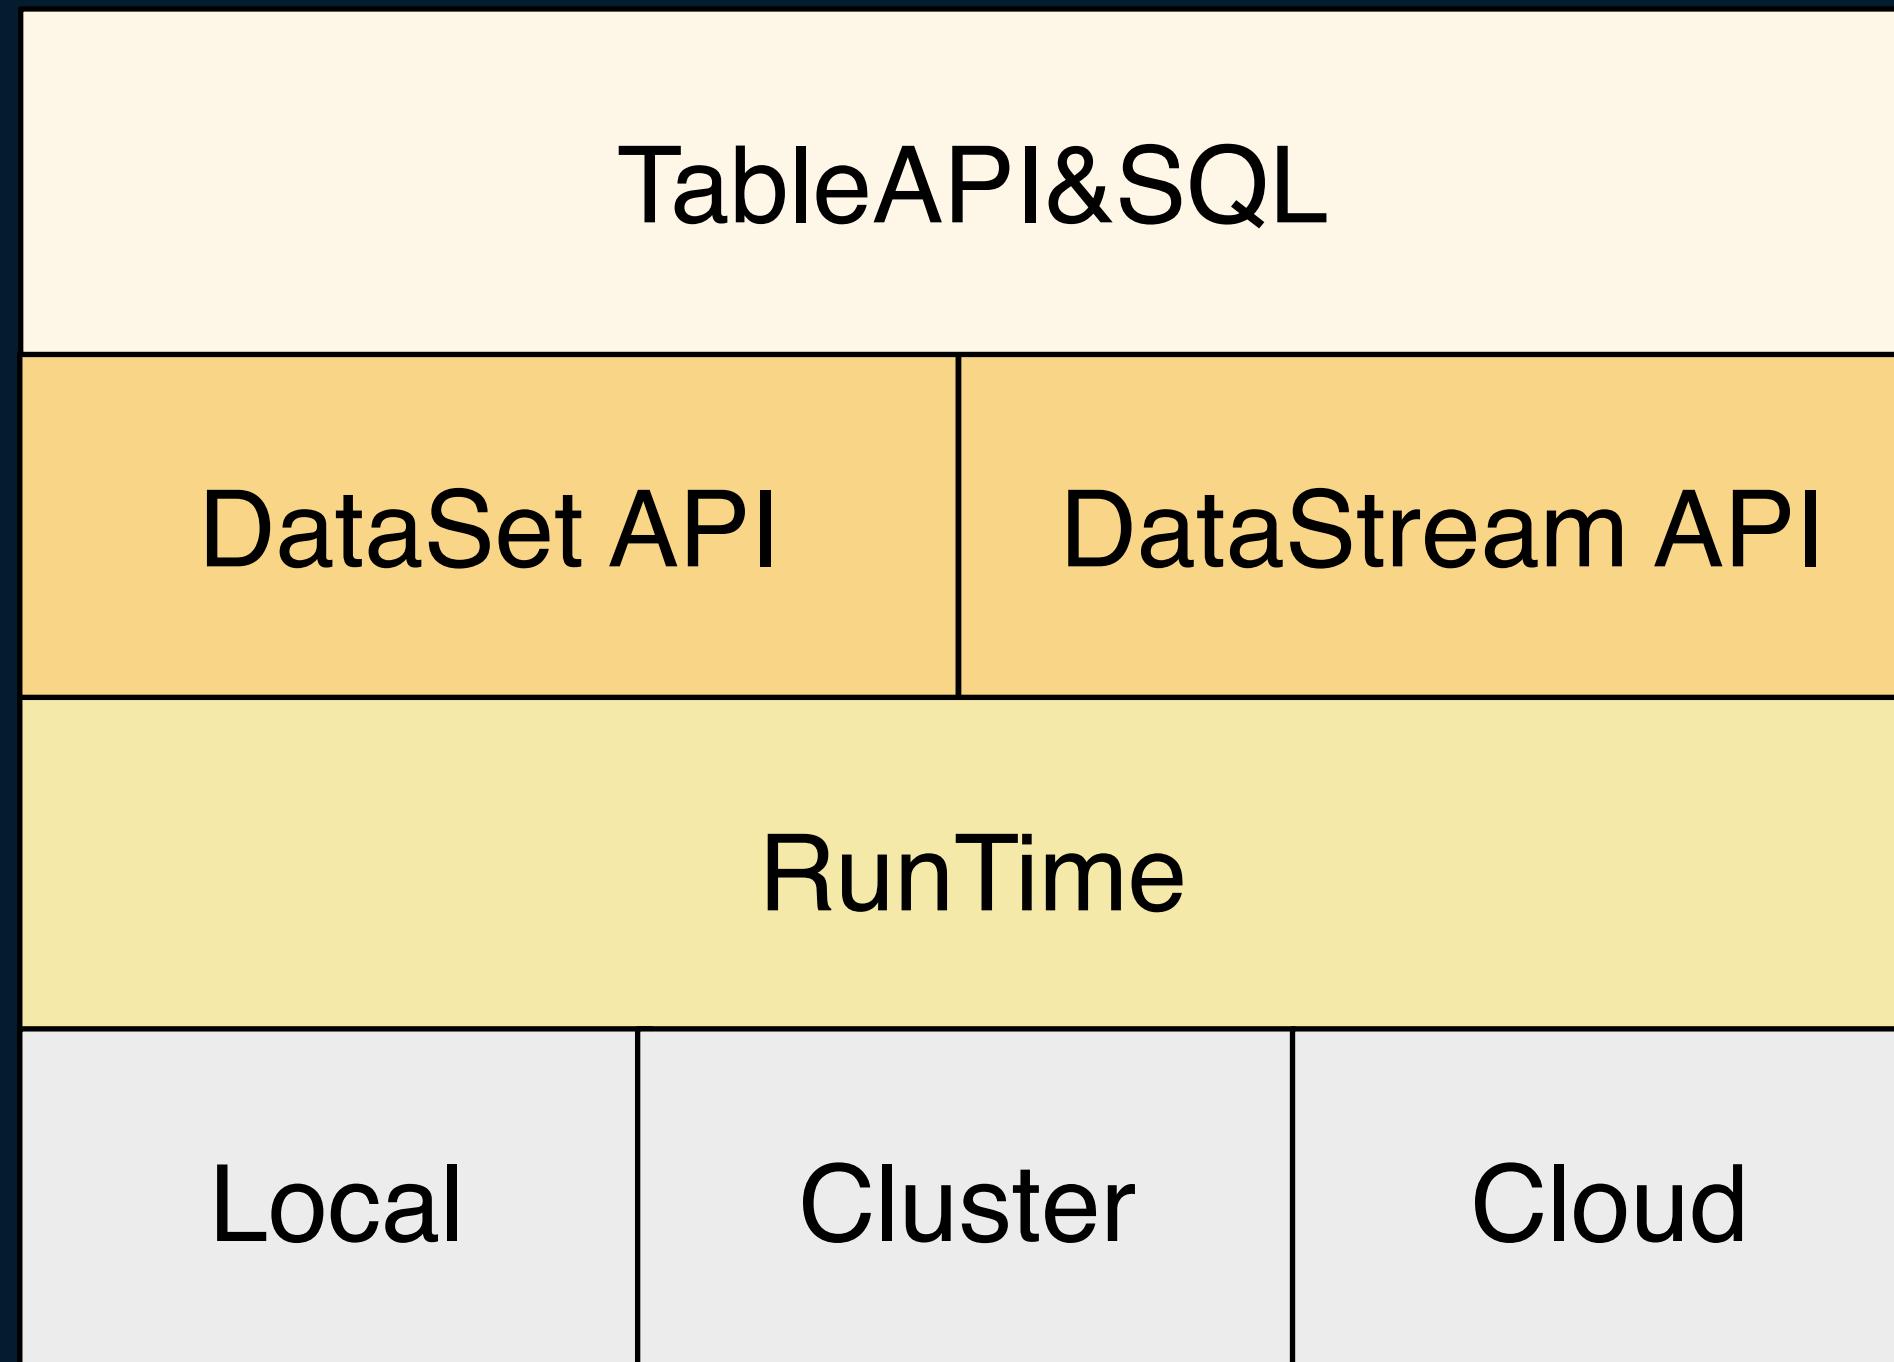


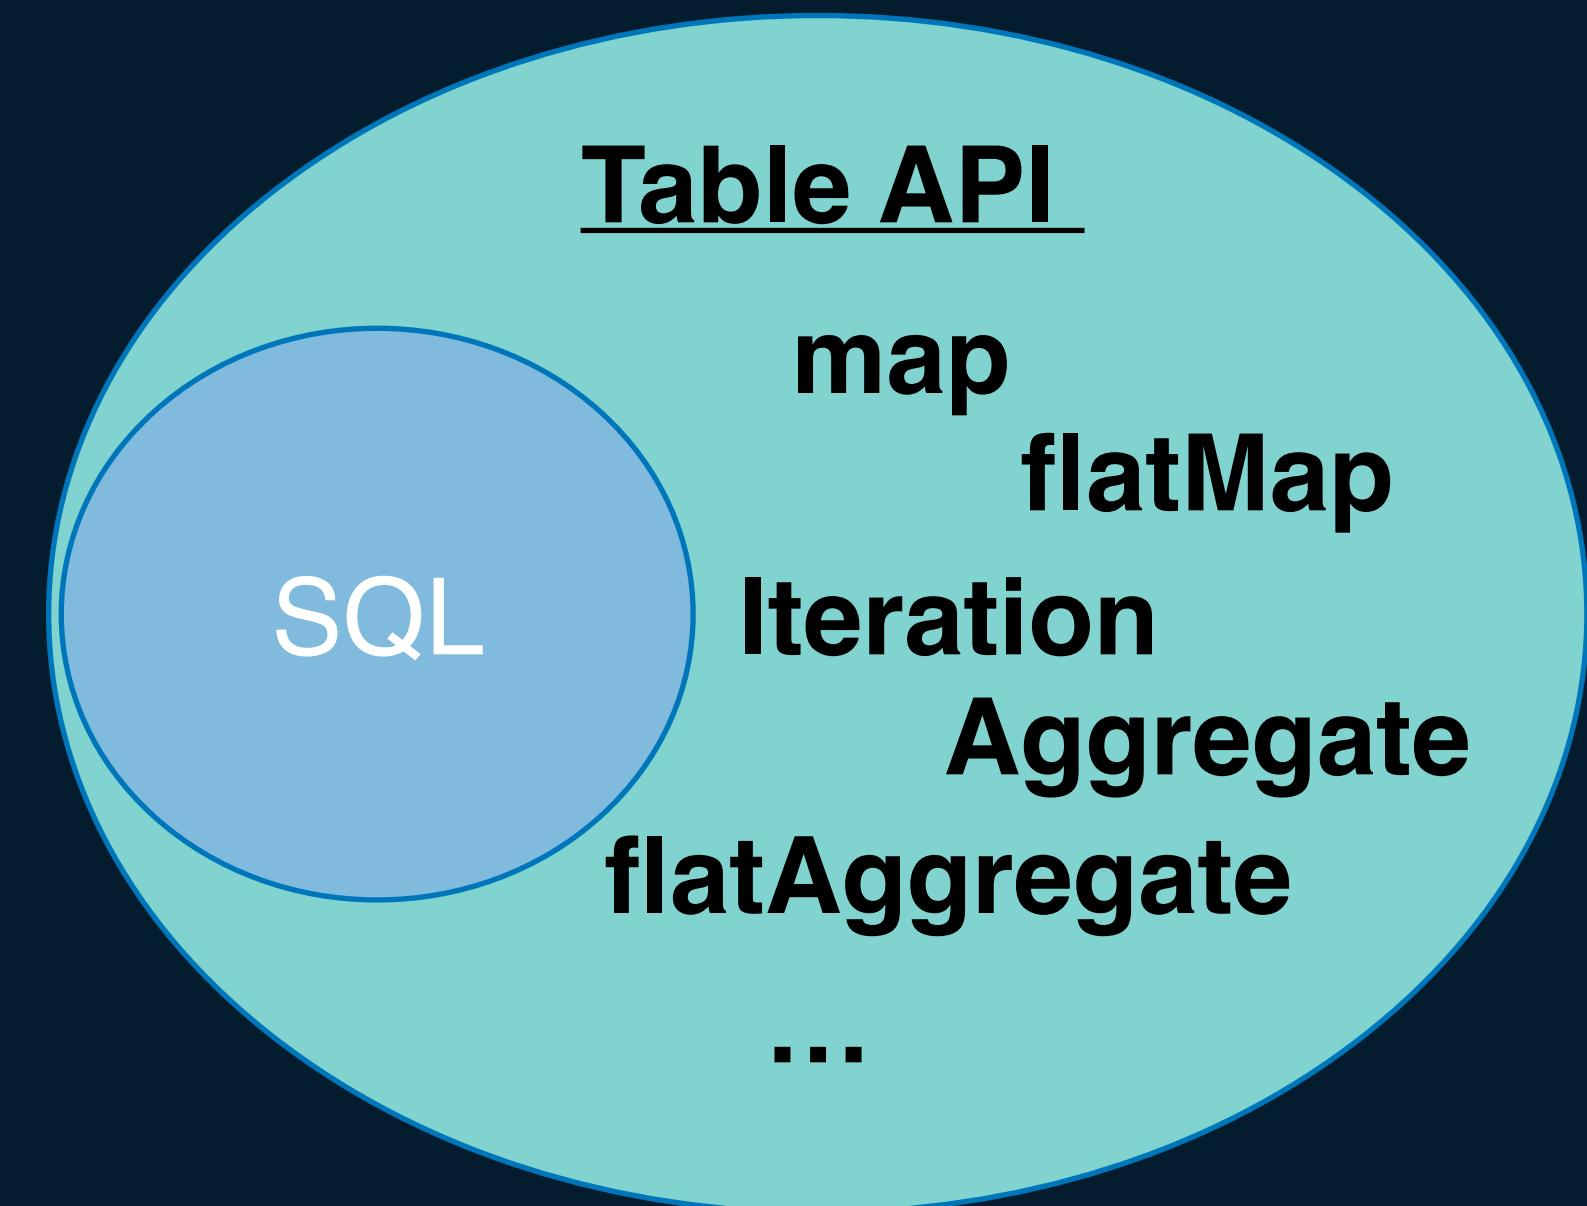
table
.groupBy('region')
.select('region, COUNT(1)')

For Bath and Streaming
支持流批两种运行模式

Table API is a super set of the SQL

Table API = SQL + ... Specially designed by Apache Flink

	TableAPI	SQL	e.g.
Stream and batch unified 流批统一	Y	Y	SELECT/AGG/ WINDOW etc.
Functional scalability 功能扩展性	Y	N	FlatAGG/Iteration/ Column operations etc.
Expressive extensibility 表达方式扩展性	Y	N	map/flatMap/ Row.flatten()/minus/ intersect etc.
Compile check 编译检查	Y	N	IDE Java/Scala



Summarize what is tableAPI

▲ Stream and batch unified
流批统一

▲ With optimization
可被优化

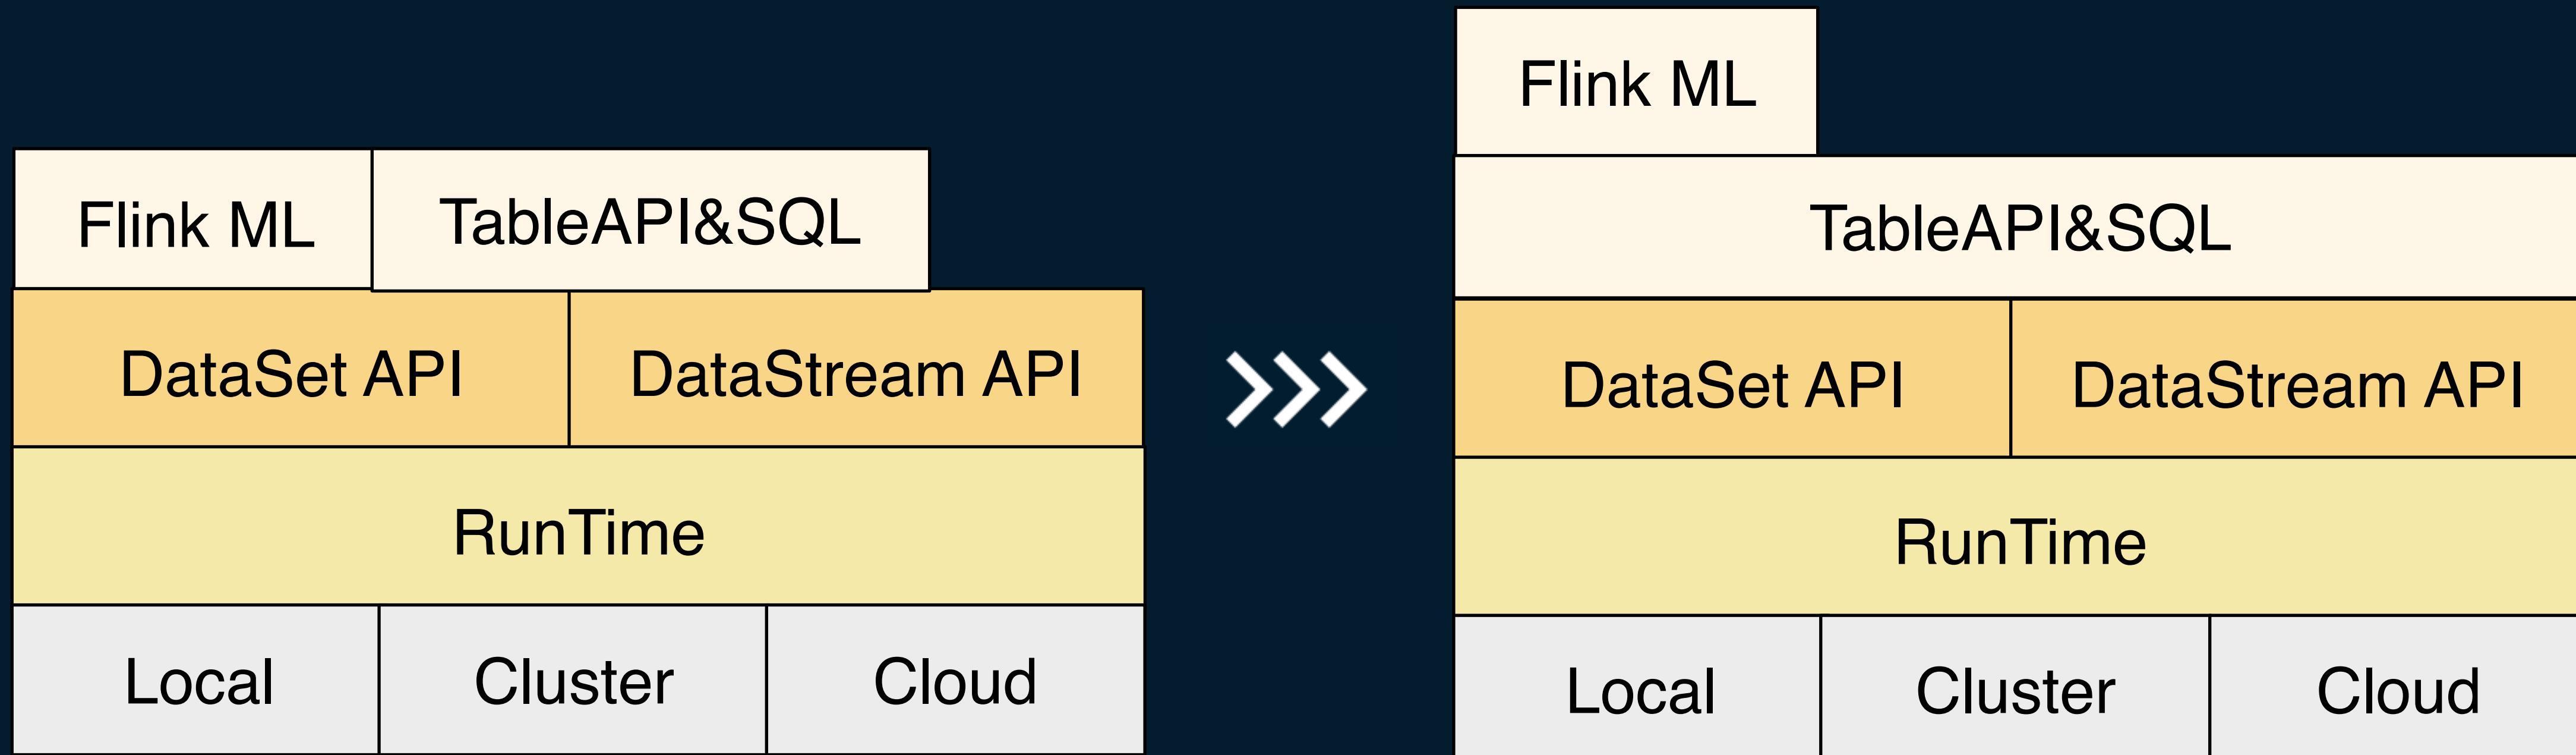
▲ Concise and easy to use
简洁易用



So, How about Flink ML on tableAPI?
那么可以基于TableAPI开发Flink ML吗?



TableAPI may unify the implementation of Flink ML



What are the requirements of Flink ML for the Table API?
Flink 机器学习对TableAPI有哪些需求呢?

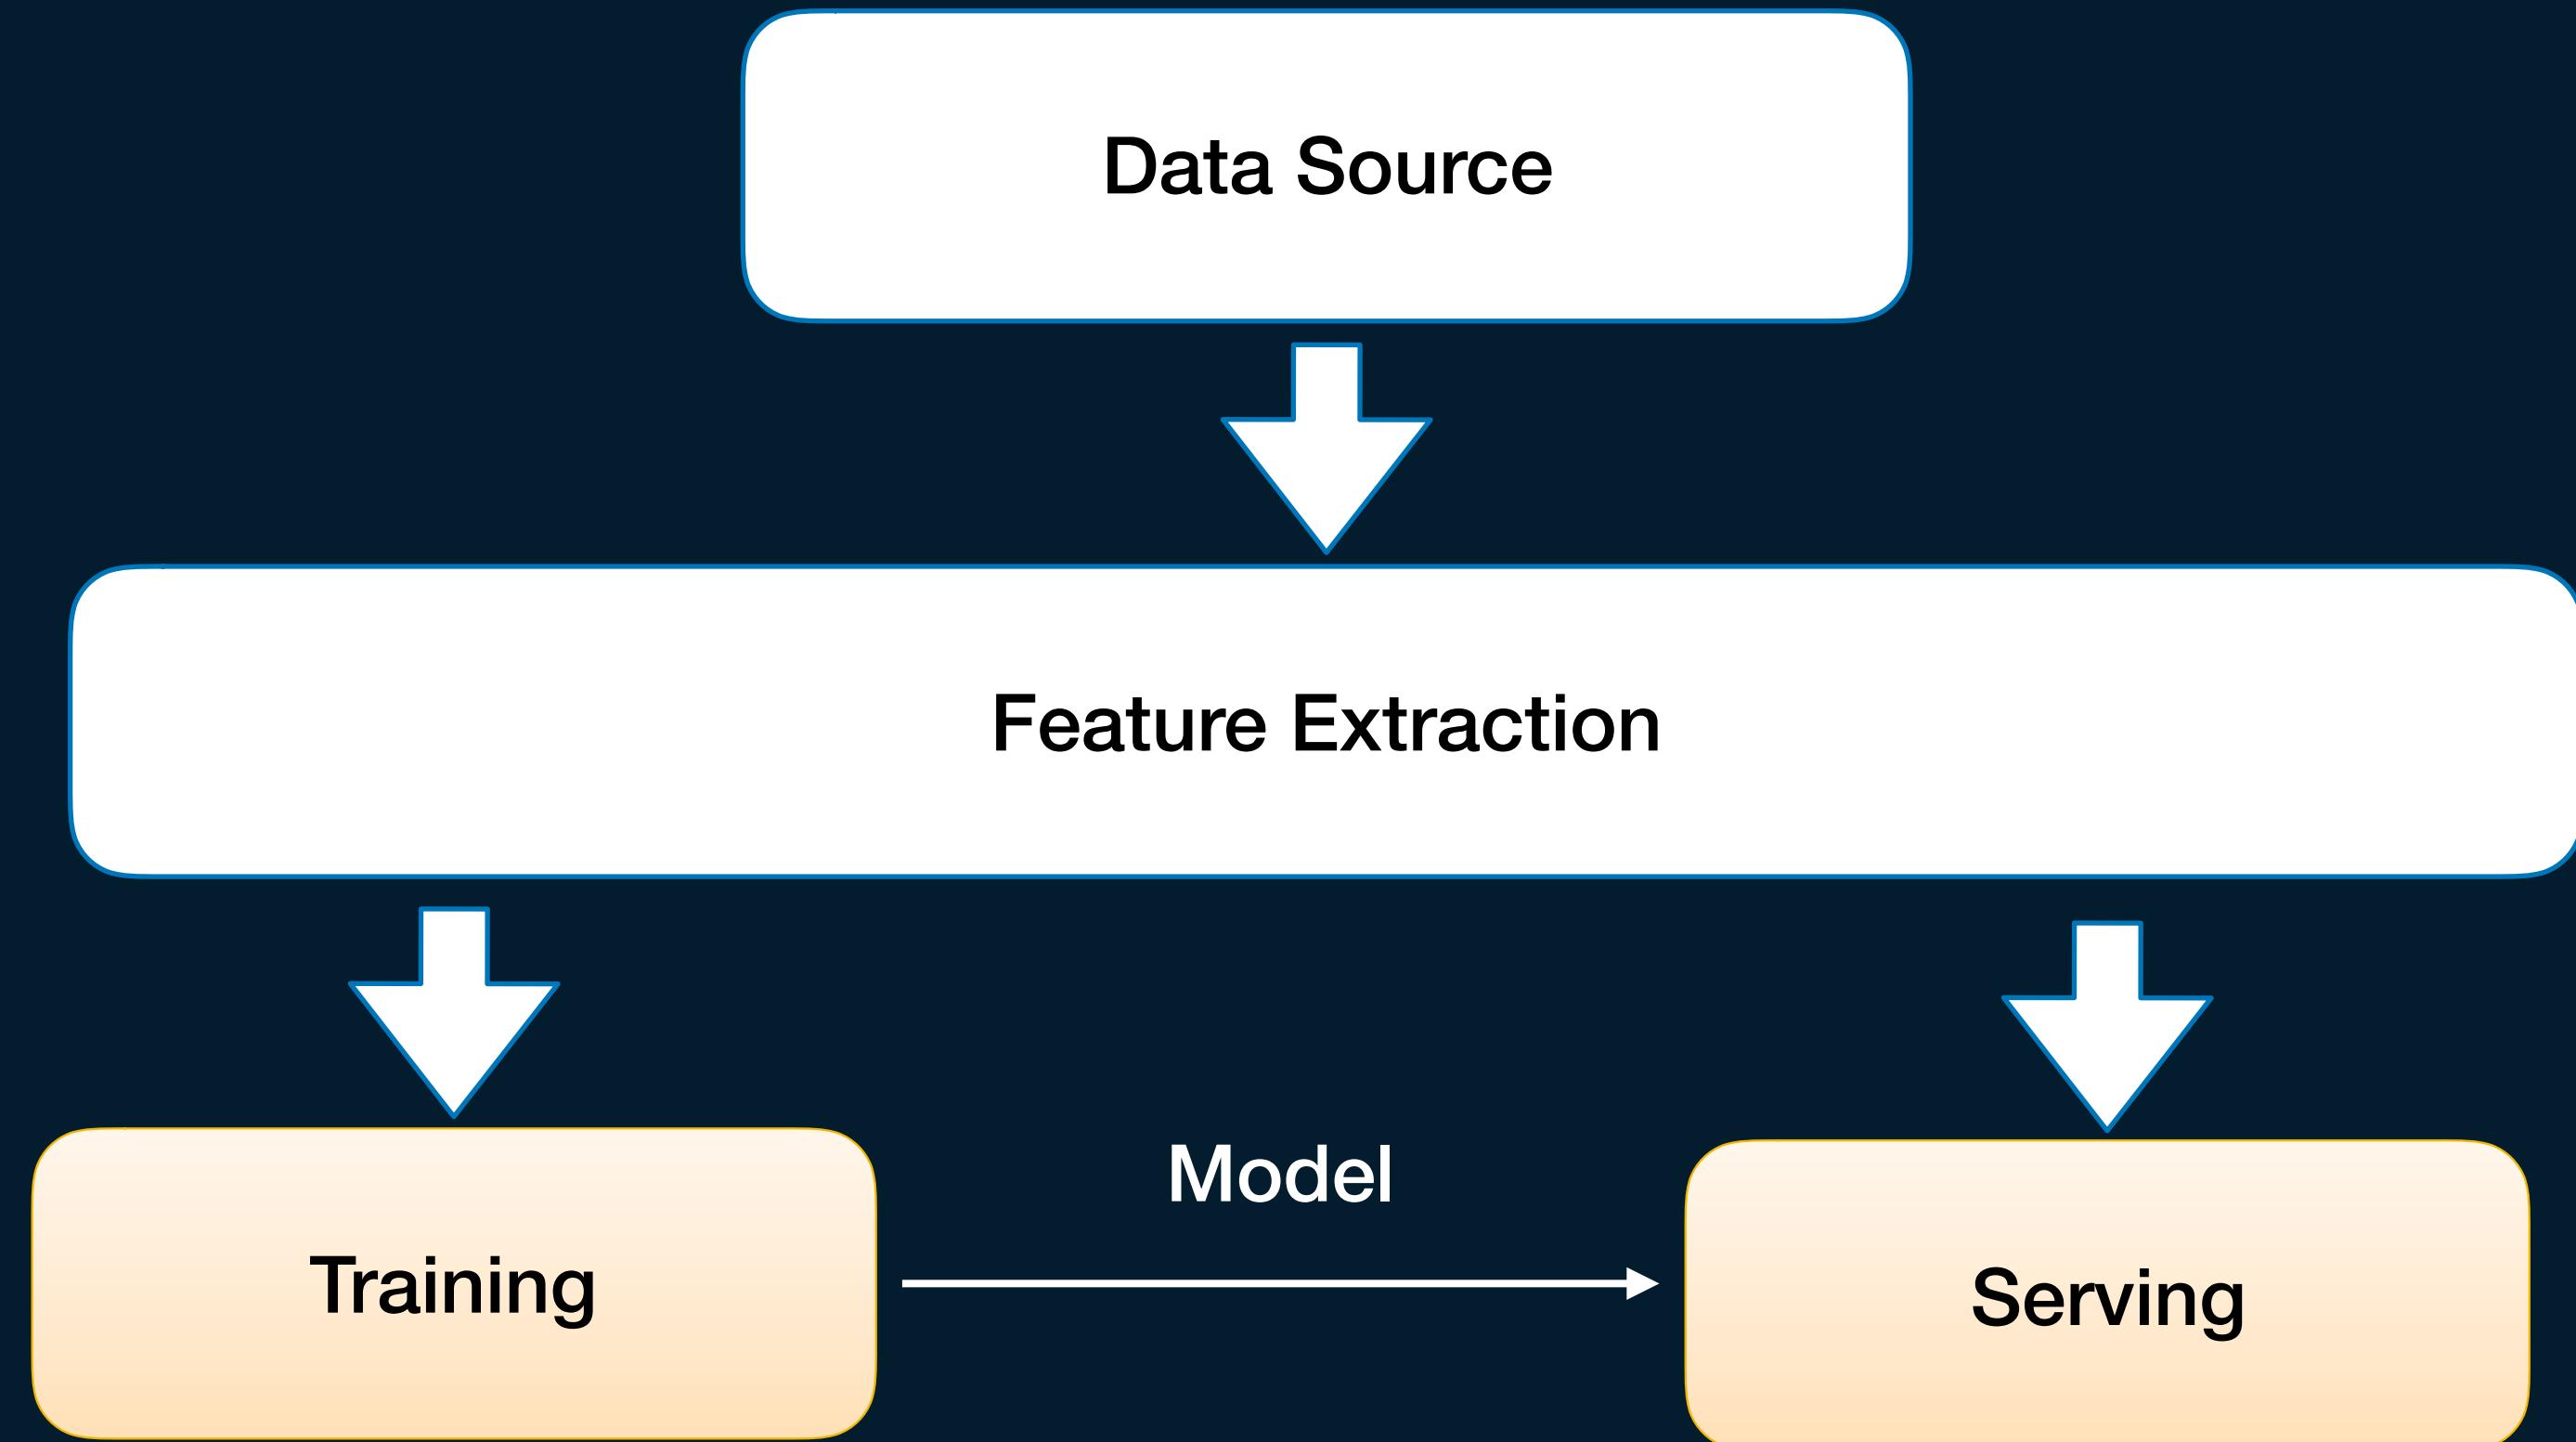
AGENDA

提要



- ▲ **Brief introduction to Flink Table API**
Table API 简介
- ▲ **The API requirements from Machine Learning Algorithms**
机器学习算法对 API 的核心需求
- ▲ **The enhancement to Flink Table API**
Flink Table API 的扩展
- ▲ **Algorithm examples based on Flink Table API**
基于 Flink Table API 的算法实现

Core API requirements from ML

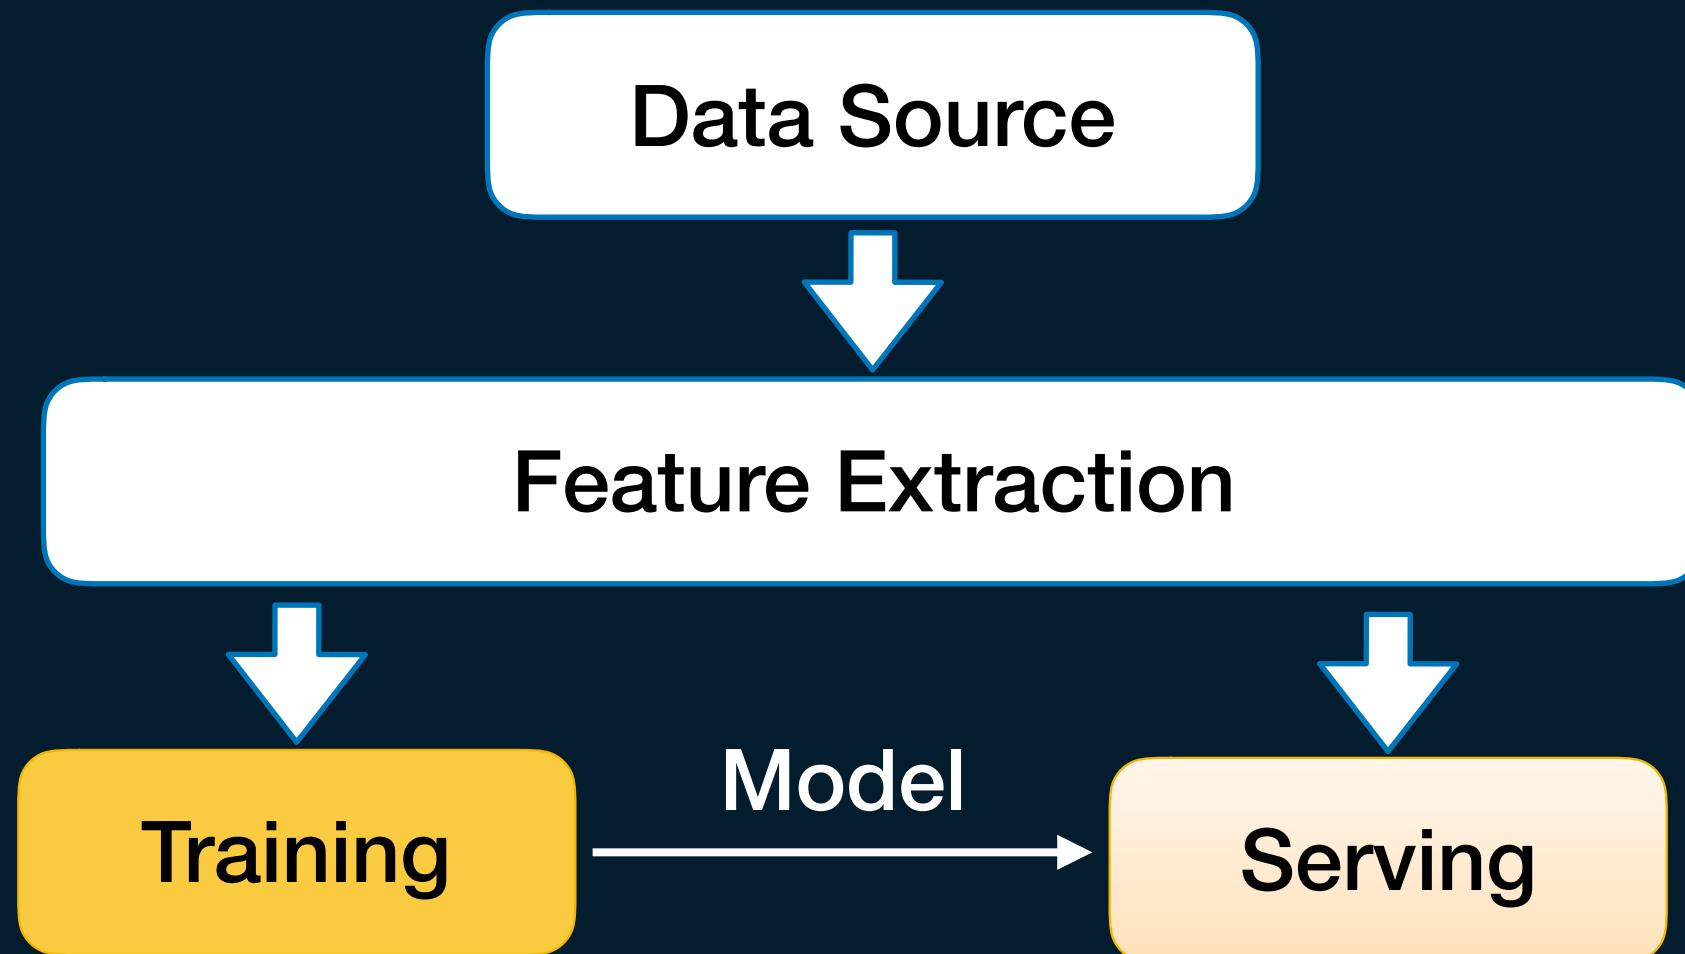


Model Training

Batch processing in most cases
主要为批计算

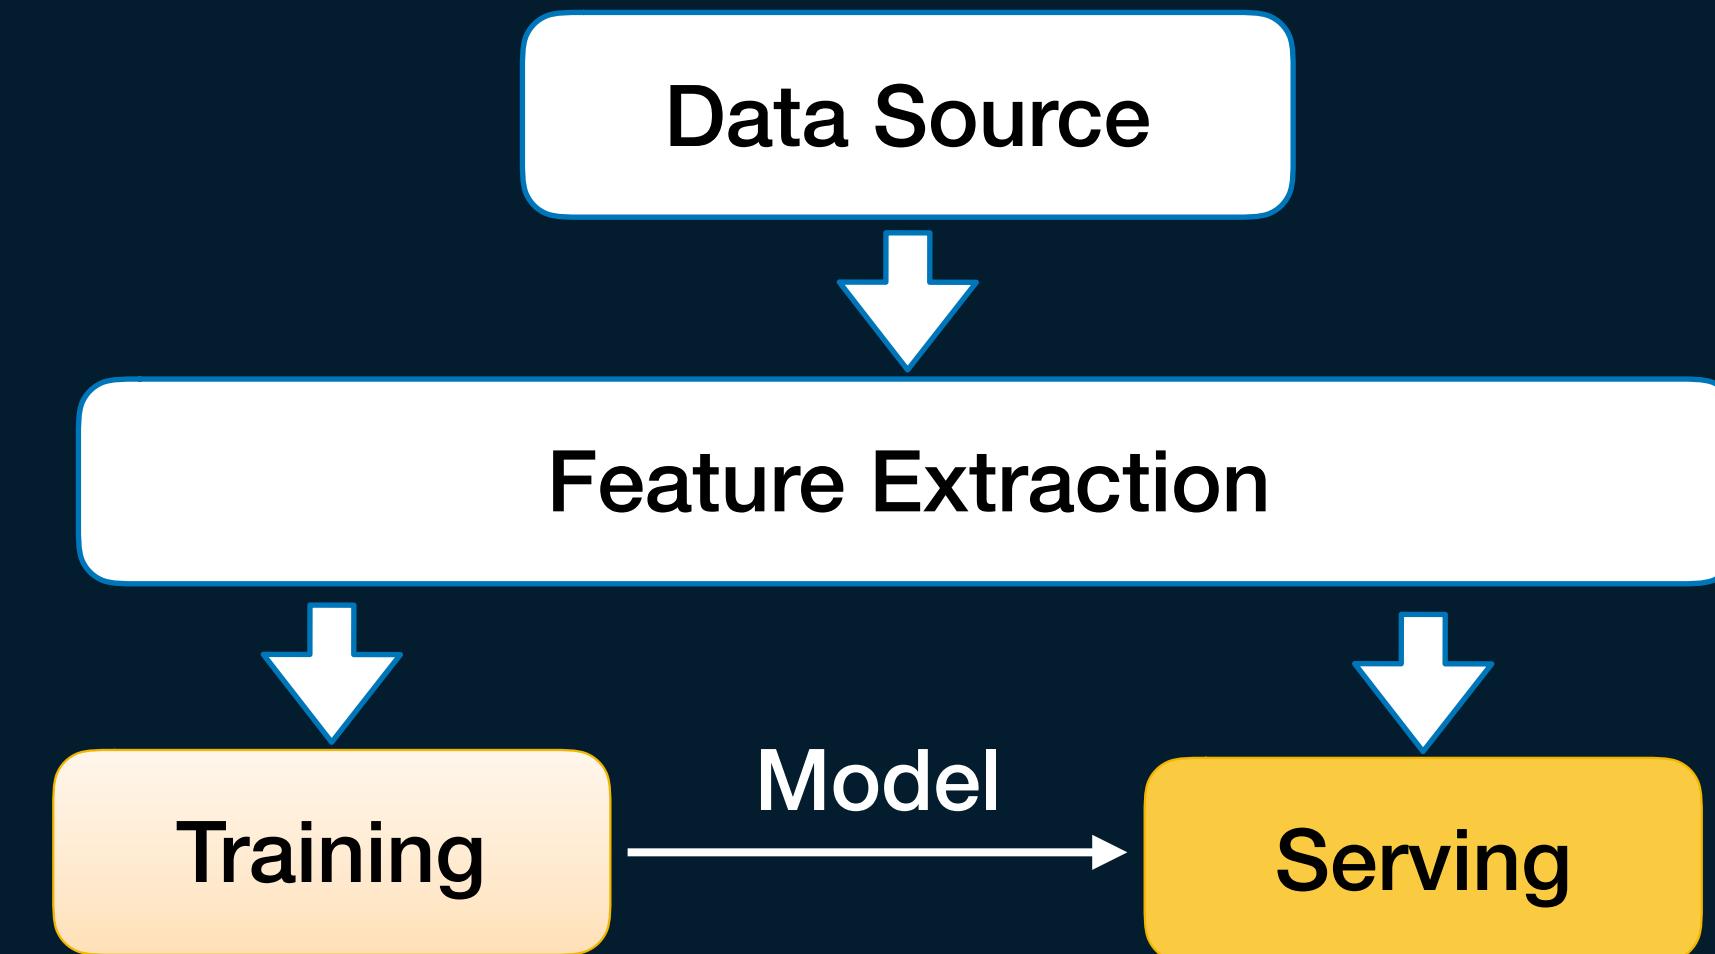
Ad-hoc algorithm experiments
算法实验

Iterate until convergence
迭代收敛



Model Serving

- 批计算和流计算
Batch processing and stream processing
- 模型的动态更新部署
Dynamic deployment of the model



From Scenarios to API Requirements



算法实验

Ad-hoc Algorithm Experiments

迭代收敛

Iterate until converge

批计算+流计算

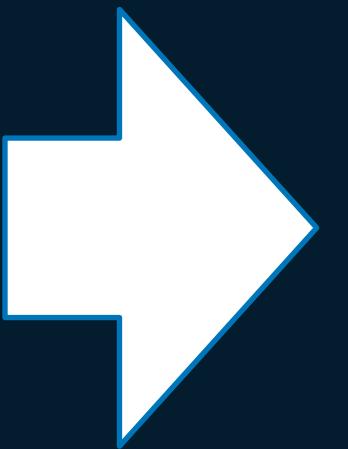
Batch + Stream processing

行计算

Row-based processing

动态模型更新部署

Dynamic model update and deployment



交互式编程

Interactive programming

迭代计算

Iterative processing

批流统一的 API

Unified API for batch and stream processing

基于整行的计算

Row-based API



From Scenarios to API Requirements



算法实验

Ad-hoc Algorithm Experiments

迭代收敛

Iterate until converge

批计算+流计算

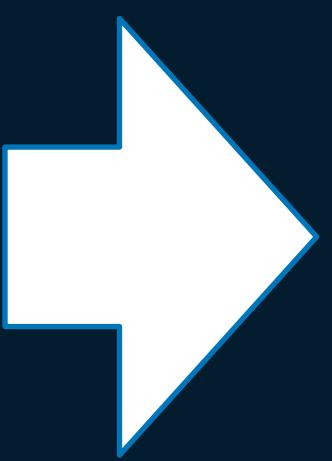
Batch + Stream processing

行计算

Row-based processing

动态模型更新部署

Dynamic model update and deployment



交互式编程 (FLINK-11199)

Interactive programming (FLINK-11199)

迭代计算 (coming soon)

Iterative processing

批流统一的 API

Unified API for batch and stream processing

基于整行的计算 (FLIP-29)

Row-based API



AGENDA

提要



Apache Flink

- ▲ **Brief introduction to Flink Table API**
Table API 简介
- ▲ **The API requirements from Machine Learning Algorithms**
机器学习算法对API的核心需求
- ▲ **The enhancement to Flink Table API**
Flink Table API 的扩展
- ▲ **Algorithm examples based on Flink Table API**
基于 Flink Table API 的算法实现

AGENDA

提要



Apache Flink

- ▲ The enhancement to Flink Table API
Flink Table API 的扩展
- ▲ 基于整行的计算及聚合功能
Row-based processing and aggregation function
- ▲ 交互式编程
Interactive programming
- ▲ 迭代计算
Iterative processing

Map Operator in Table API

Method signature 方法签名

```
def map(scalaFunction: Expression): Table
```

Usage 用法

```
val res = tab
  .map(fun('e)).as('a, 'b, 'c)
  .select('a, 'c)
```

Benefit 好处

`table.select(udf1(), udf2(), udf3()....)`
 vs
`table.map(udf())`

```
class MyMap extends ScalarFunction {
  var param : String = ""

  override def open(context: FunctionContext): Unit
    = param = context.getJobParameter("paramKey","")

  def eval([user defined inputs]): Row = {
    val result = new Row(3)
    // Business processing based on data and parameters
    // 根据数据和参数进行业务处理
    result
  }

  override def getResultType(signature:
Array[Class[_]]): TypeInformation[_] = {
    Types.ROW(Types.STRING, Types.INT, Types.LONG)
  }
}
```

INPUT(Row)	OUTPUT(Row)
1	1

FlatMap Operator in Table API

Method signature

方法签名

```
def flatMap(tableFunction: Expression): Table
```

Usage

用法

```
val res = tab
  .flatMap(fun('e,'f)).as('name, 'age)
  .select('name, 'age)
```

Benefit

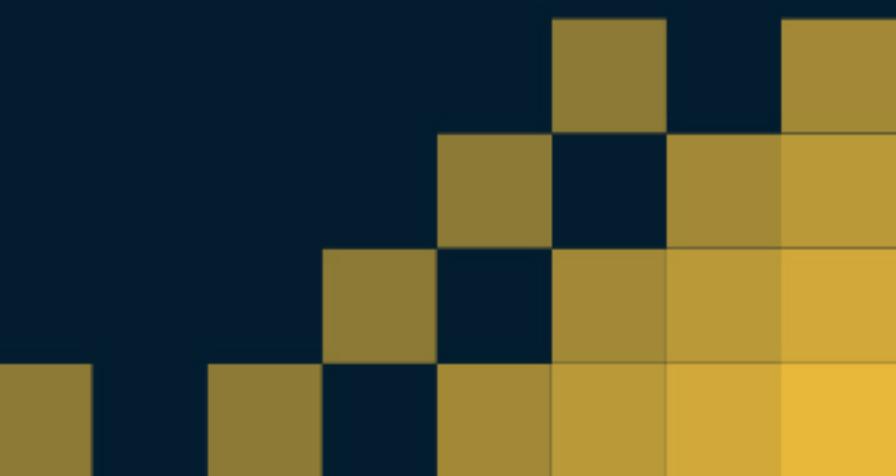
好处

table.join(udtf) VS table.flatMap(udtf())

```
case class User(name: String, age: Int)

class MyFlatMap extends TableFunction[User] {
  def eval([user defined inputs]): Unit = {
    for(...){
      collect(User(name, age))
    }
  }
}
```

INPUT(Row)	OUTPUT(Row)
1	N(N>=0)



Aggregate Operator in Table API

Method signature

方法签名

```
def aggregate(aggregateFunction: Expression): AggregatedTable
```

```
class AggregatedTable(table: Table, groupKeys: Seq[Expression], aggFunction: Expression)
```

Usage

用法

```
val res = groupedTab
    .groupBy('a)
    .aggregate(agg('e,'f) as ('a, 'b, 'c))
    .select('a, 'c)
```

Benefit

好处

table.select(agg1(), agg2(), agg3()....)
VS
table.aggregate(agg())

```
class CountAccumulator extends JTuple1[Long] {
  f0 = 0L //count
}

class CountAgg extends AggregateFunction[JLong, CountAccumulator] {
  def accumulate(acc: CountAccumulator): Unit = {
    acc.f0 += 1L
  }

  override def getValue(acc: CountAccumulator): JLong = {
    acc.f0
  }
  ... retract()/merge()
}
```

INPUT(Row)	OUTPUT(Row)
N(N>=0)	1

FlatAggregate Operator in Table API

Method signature

方法签名

```
def flatAggregate(tableAggregateFunction: Expression): GroupedFlatAggregateTable
```

```
class GroupedFlatAggTable(table: Table, groupKey: Seq[Expression], tableAggFun: Expression)
```

Usage

用法

```
val res = groupedTab
    .groupBy('a)
    .flatAggregate(
        flatAgg('e,'f) as ('a, 'b, 'c))
    .select('a, 'c)
```

Benefit

好处

New features on the TableAPI
TableAPI 上面的新增功能

```
class TopNAcc {
    var data: MapView[JInt, JLong] = _ // (rank -> value)
    ...
}

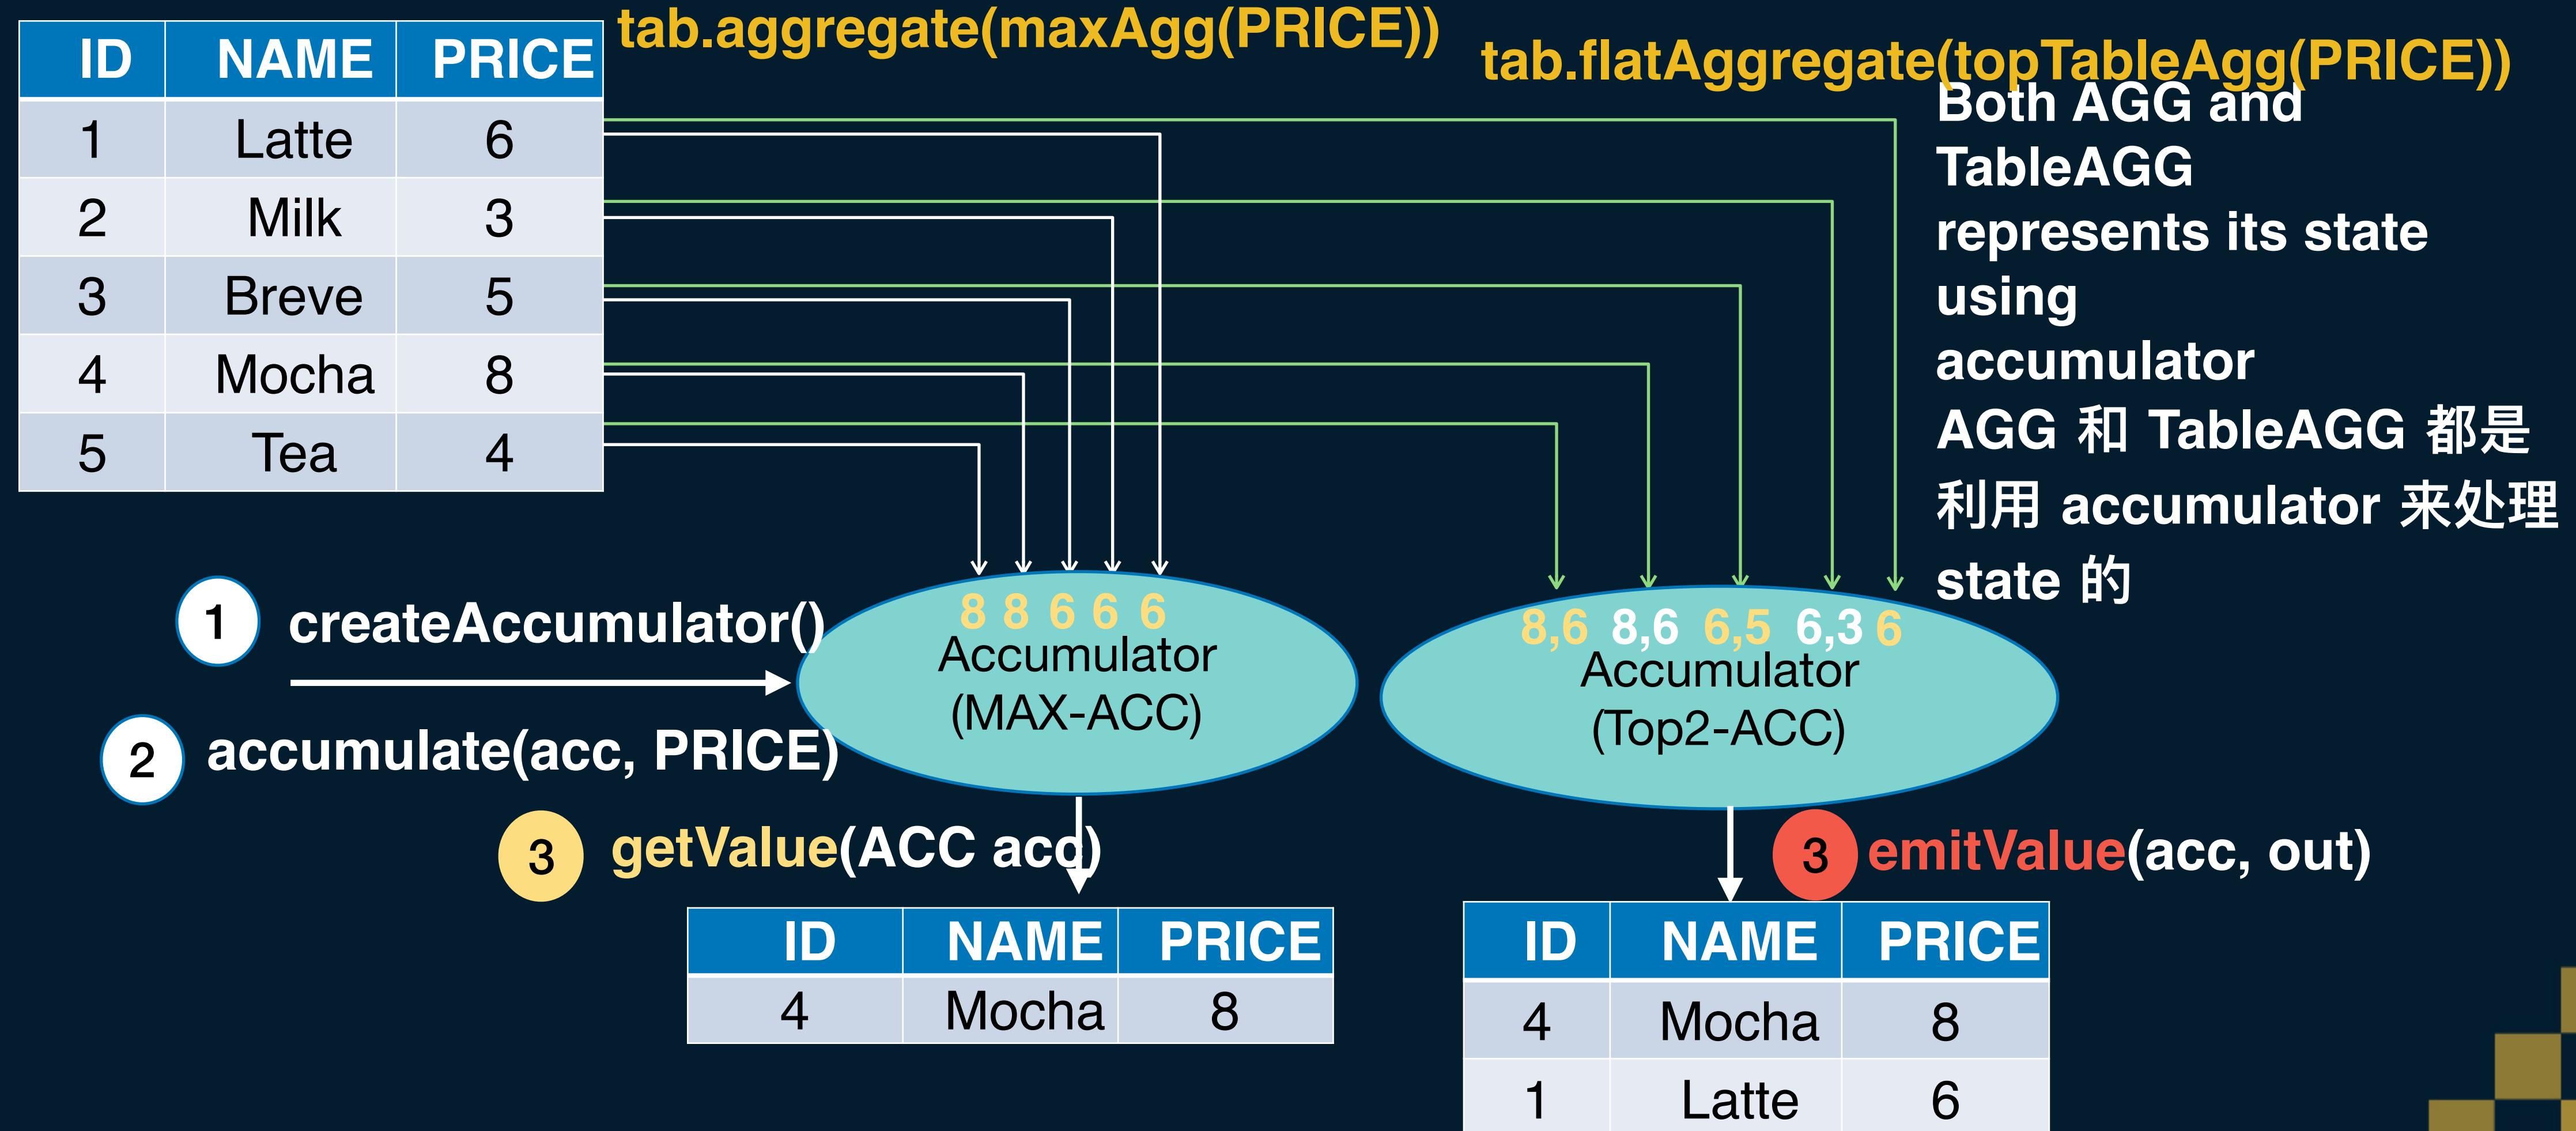
class TopN(n: Int) extends TableAggregateFunction[(Int, Long),
TopNAccum] {

    def accumulate(acc: TopNAcc, [user defined inputs]): Unit = {
        ...
    }
    def emitValue(acc: TopNAcc, out: Collector[(Int, Long)]): Unit = {
        ...
    }
    ...
    ... retract/merge
}
```

INPUT(Row)	OUTPUT(Row)
N(N>=0)	M(M>=0)

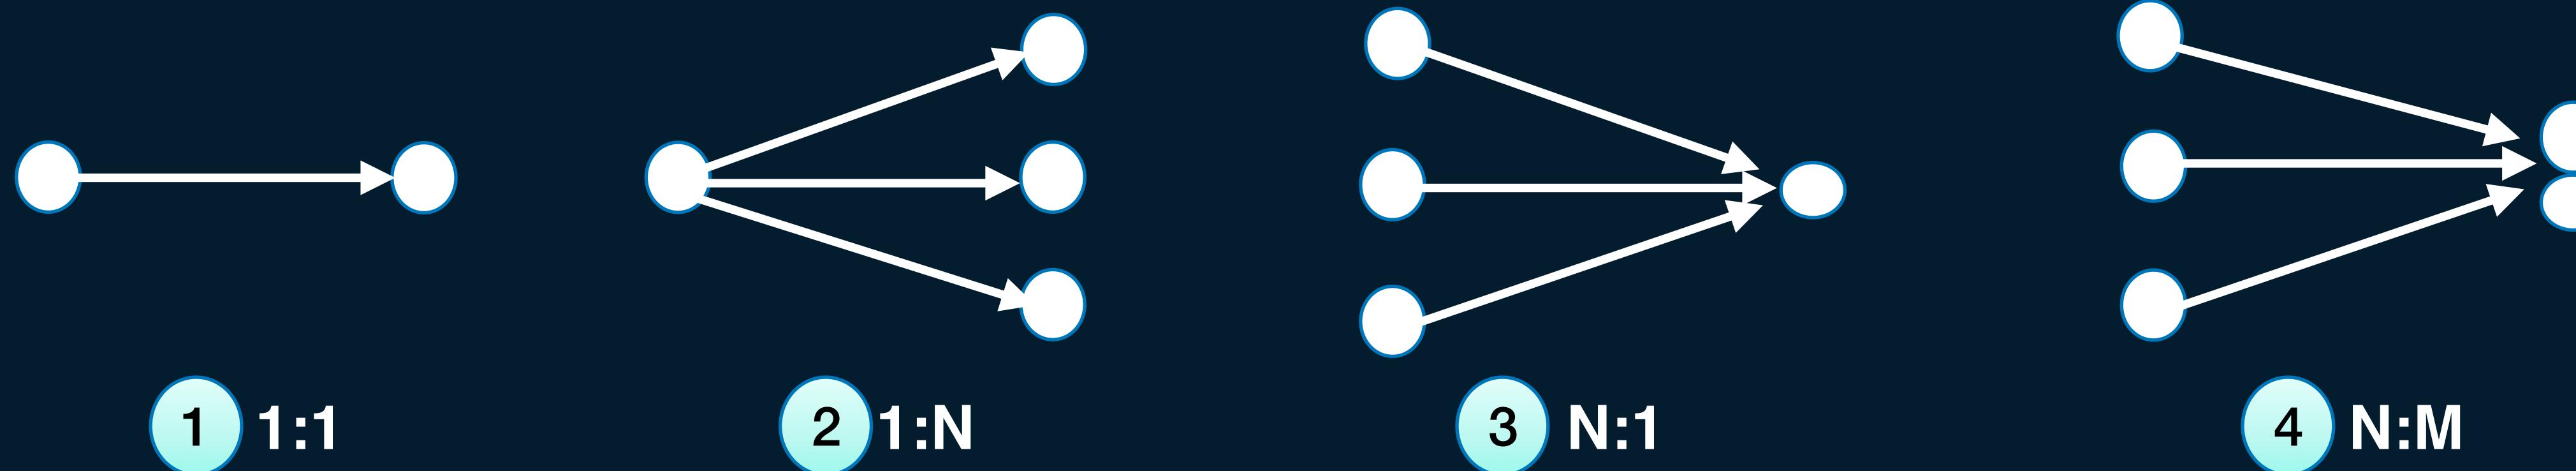
Aggregate VS FlatAggregate

Scenes using Max and Top2 compare the difference between Aggregate and FlatAggregate
使用 Max 和 TopN 的场景比较 Aggregate 和 FlatAggregate 之间的差异



Summarize the enhancements to TableAPI

	Single Row Input 单行输入	Multiple Row Input 多行输入
Single Row Output 单行输出	ScalarFunction (select/map)	AggregateFunction (select/aggregate)
Multiple Row Output 多行输出	TableFunction (cross join/flatmap)	TableAggregateFunction (flatAggregate)



AGENDA

提要



Apache Flink

The enhancement to Flink Table API Flink Table API的扩展

基于整行的计算及聚合功能

Row-based processing and aggregation function

Interactive programming

交互式编程

Iterative processing

迭代计算

A example code snippet

{

```
  val orders = tEnv.fromCollection(data).as ('country, 'color, 'quantity)
```

```
  val smallOrders = orders.filter('quantity < 100)
```

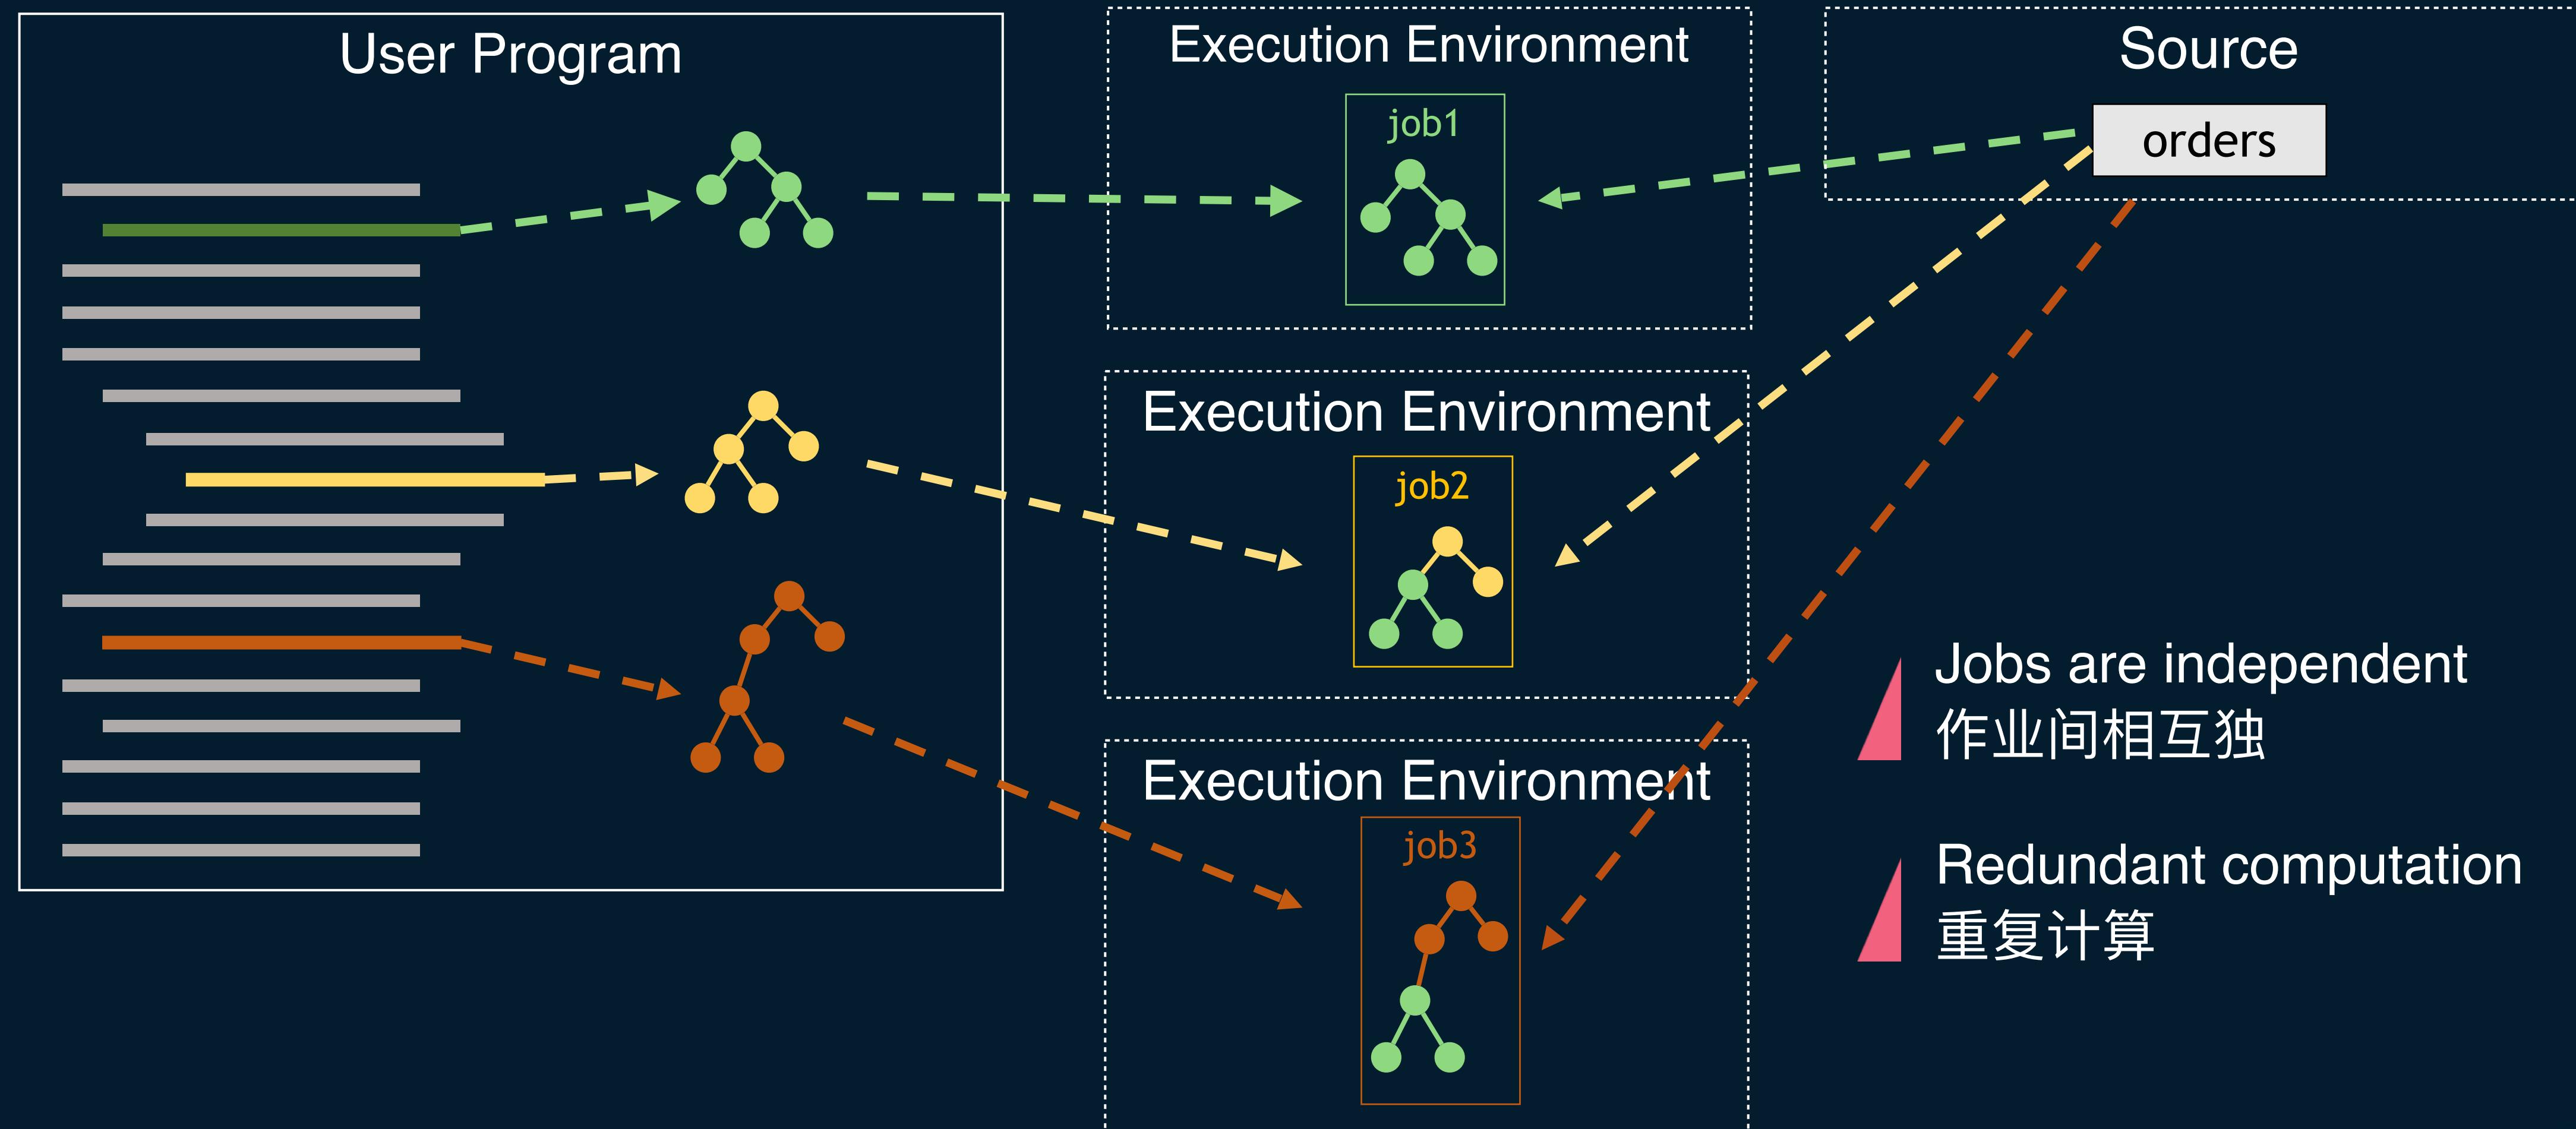
```
  val countriesOfSmallOrders = smallOrders.select('country).distinct()  
  countriesOfSmallOrders.print()
```

```
  val smallOrdersByCountry = smallOrders.groupBy('country).select('country, 'quantity.sum as 'TotalSales)  
  smallOrdersByCountry.print()
```

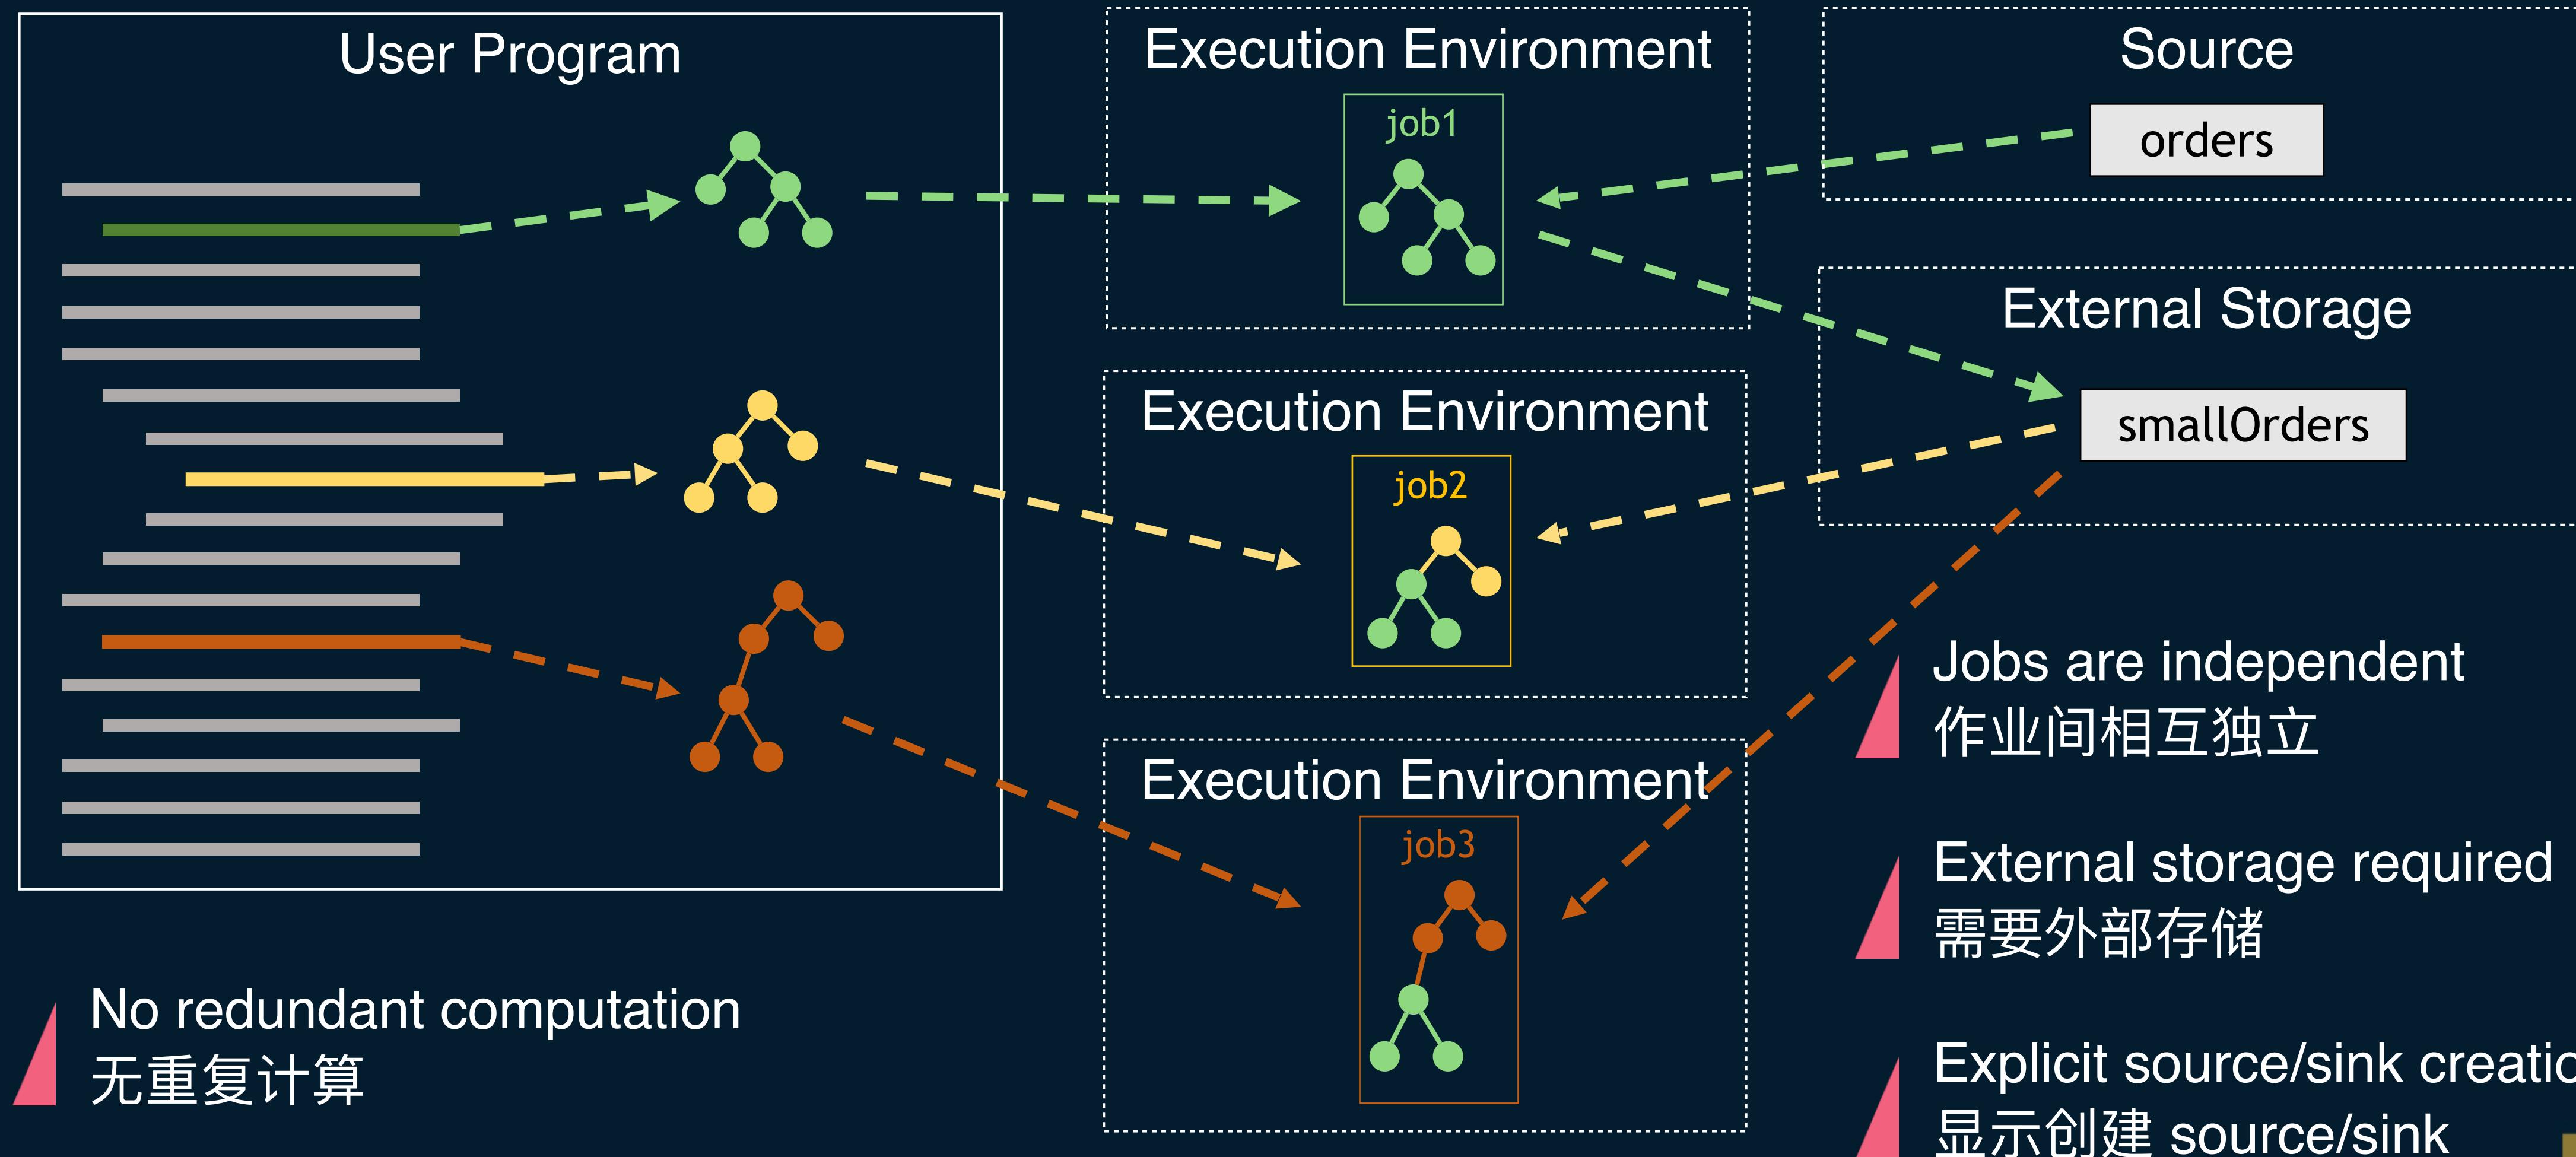
```
  val smallOrdersByColor = smallOrders.groupBy('color).select('color, 'quantity.avg as 'avg)  
  smallOrdersByColor.print()
```

}

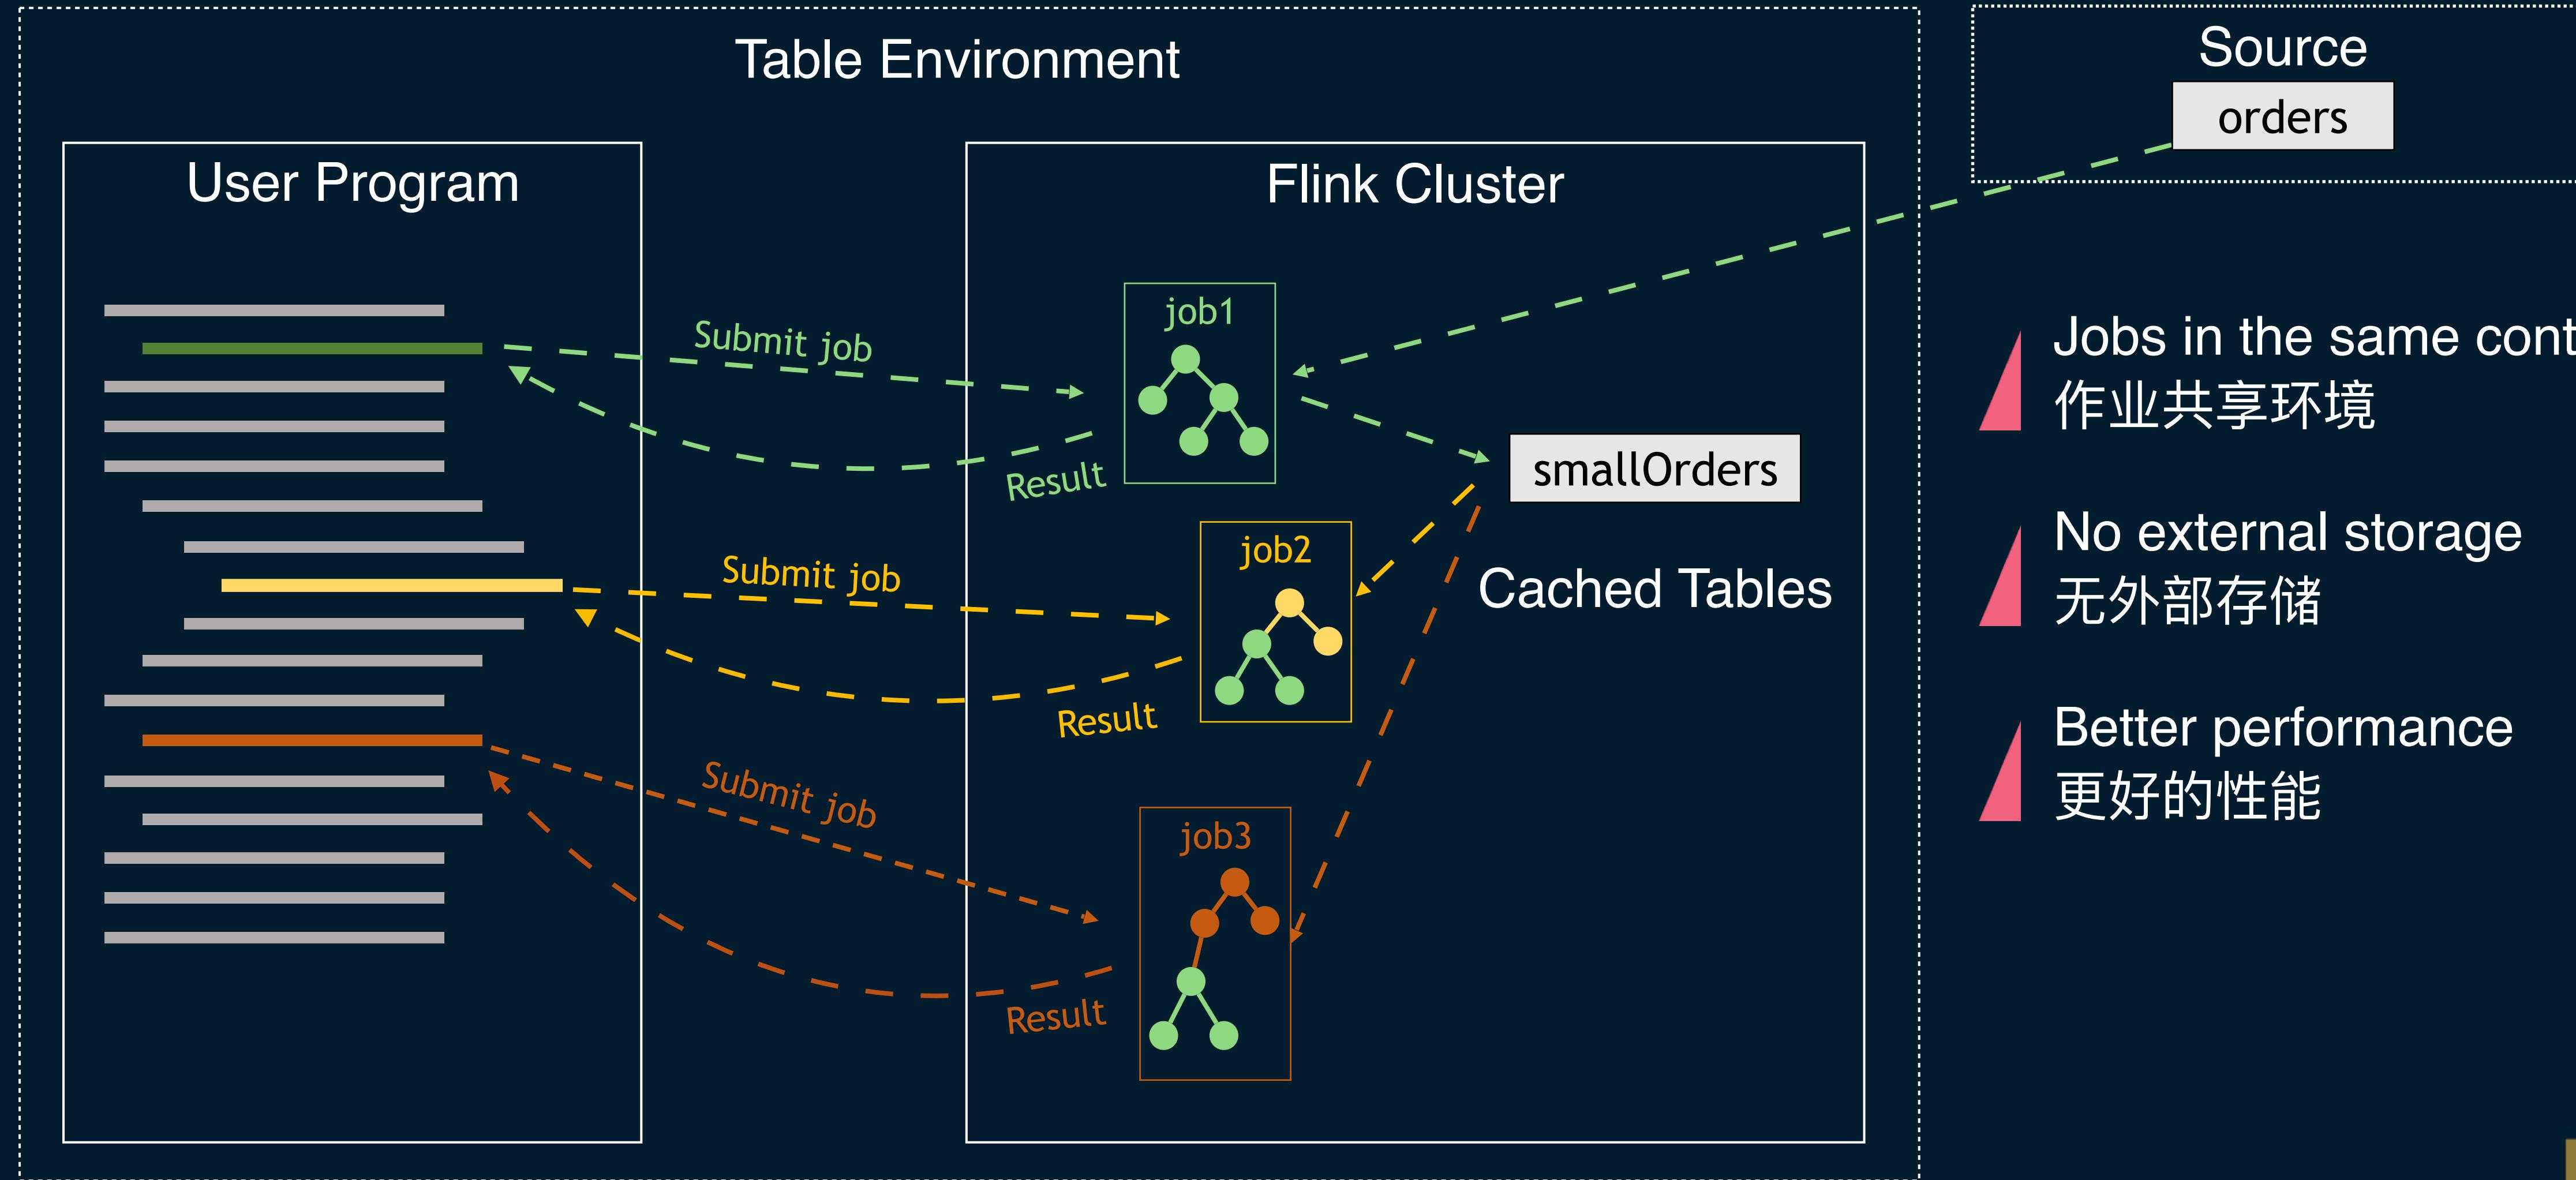
Static programming



A workaround



Interactive Programming



Interactive Programming



{

```
  val orders = tEnv.fromCollection(data).as ('country, 'color, 'quantity)
```

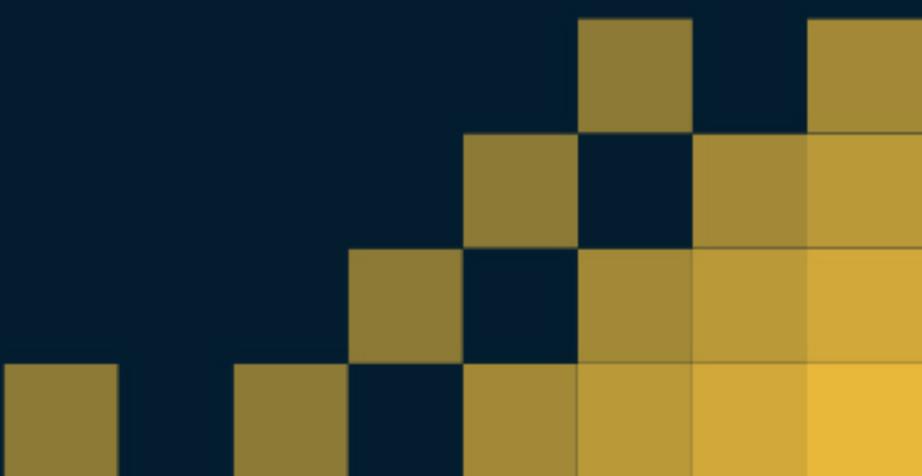
```
  val smallOrders = orders.filter('quantity < 100)  
  smallOrders.cache()
```

```
  val countriesOfSmallOrders = smallOrders.select('country).distinct()  
  countriesOfSmallOrders.print()
```

```
  val smallOrdersByCountry = smallOrders.groupBy('country).select('country, 'quantity.sum as 'TotalSales)  
  smallOrdersByCountry.print()
```

```
  val smallOrdersByColor = smallOrders.groupBy('color).select('color, 'quantity.avg as 'avg)  
  smallOrdersByColor.print()
```

}



AGENDA

提要



Apache Flink

The enhancement to Flink Table API Flink Table API 的扩展

基于整行的计算及聚合功能
Row-based processing and aggregation function

Interactive programming
交互式编程

Iterative processing
迭代计算

Iterative processing

- Execute same processing logic repeatedly
重复执行相同的计算逻辑
- Input of each round is from the output of the previous round
每次计算的变量输入来自于上次计算的结果



Components of iteration

Input variables (updated in iterations)
迭代变量 (迭代中更新)

```
{  
    val seq = Seq(1, 2, 3)
```

Step function (update input variables)
单次计算逻辑 (更新迭代变量)

```
for (i <- 1 until 100) {  
    seq.foreach(_ + 1)  
}
```

Termination condition
迭代终止条件

```
}
```



Two ways of doing iteration

Iteration using interactive programming
利用交互式编程实现

Native iteration with descriptive API
描述性API实现的原生迭代



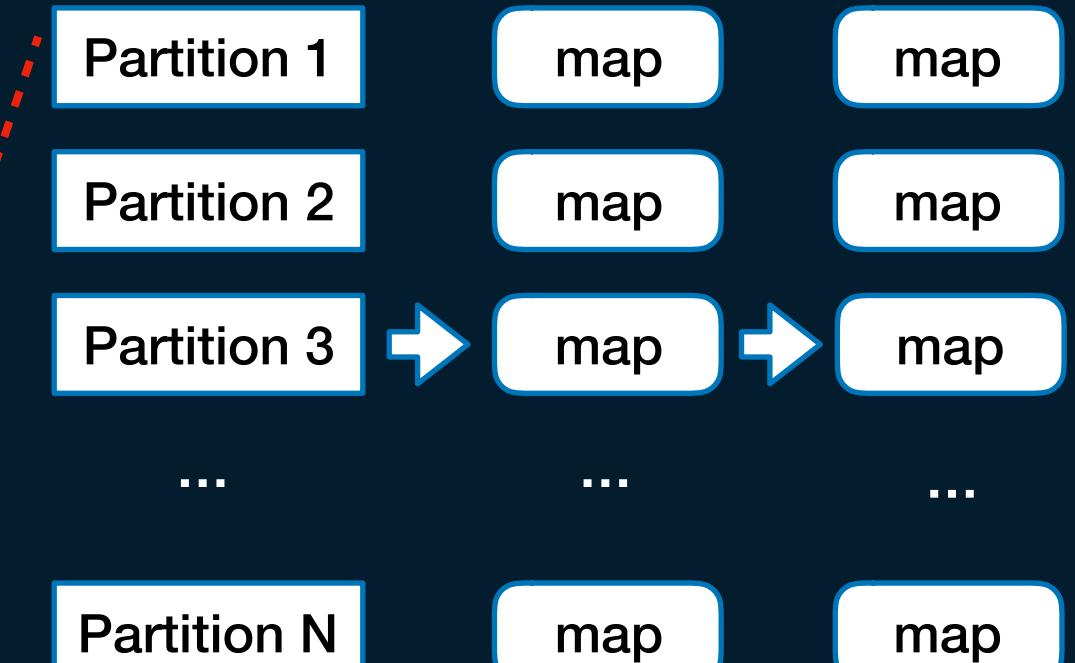
Iterative processing with interactive programming



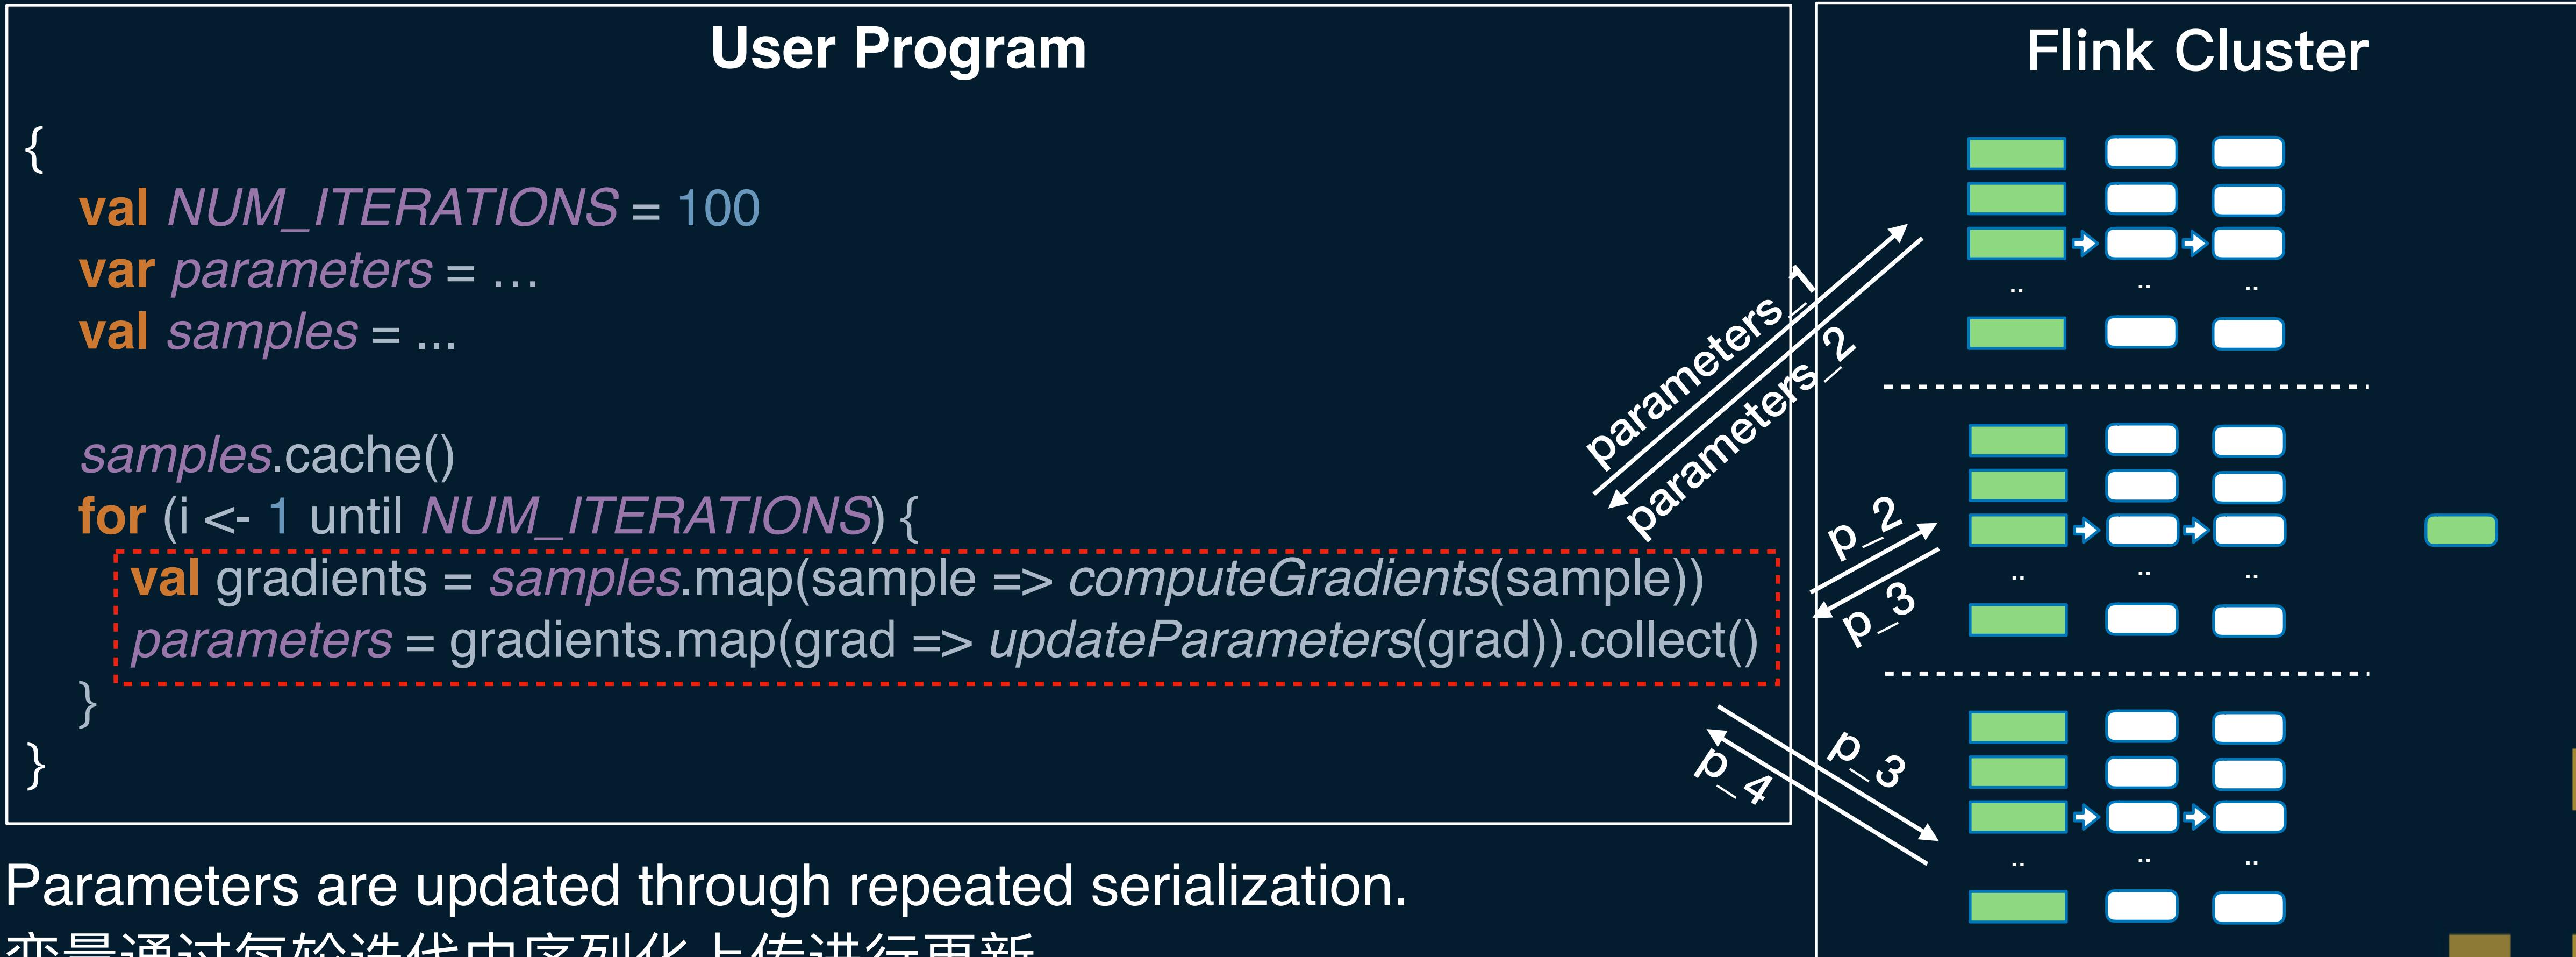
User Program

```
{  
    val NUM_ITERATIONS = 100  
    var parameters = ...  
    val samples = ...  
  
    samples.cache()  
    for(i <- 1 until NUM_ITERATIONS) {  
        val gradients = samples.map(sample => computeGradients(sample))  
        parameters = gradients.map(grad => updateParameters(grad)).collect()  
    }  
}
```

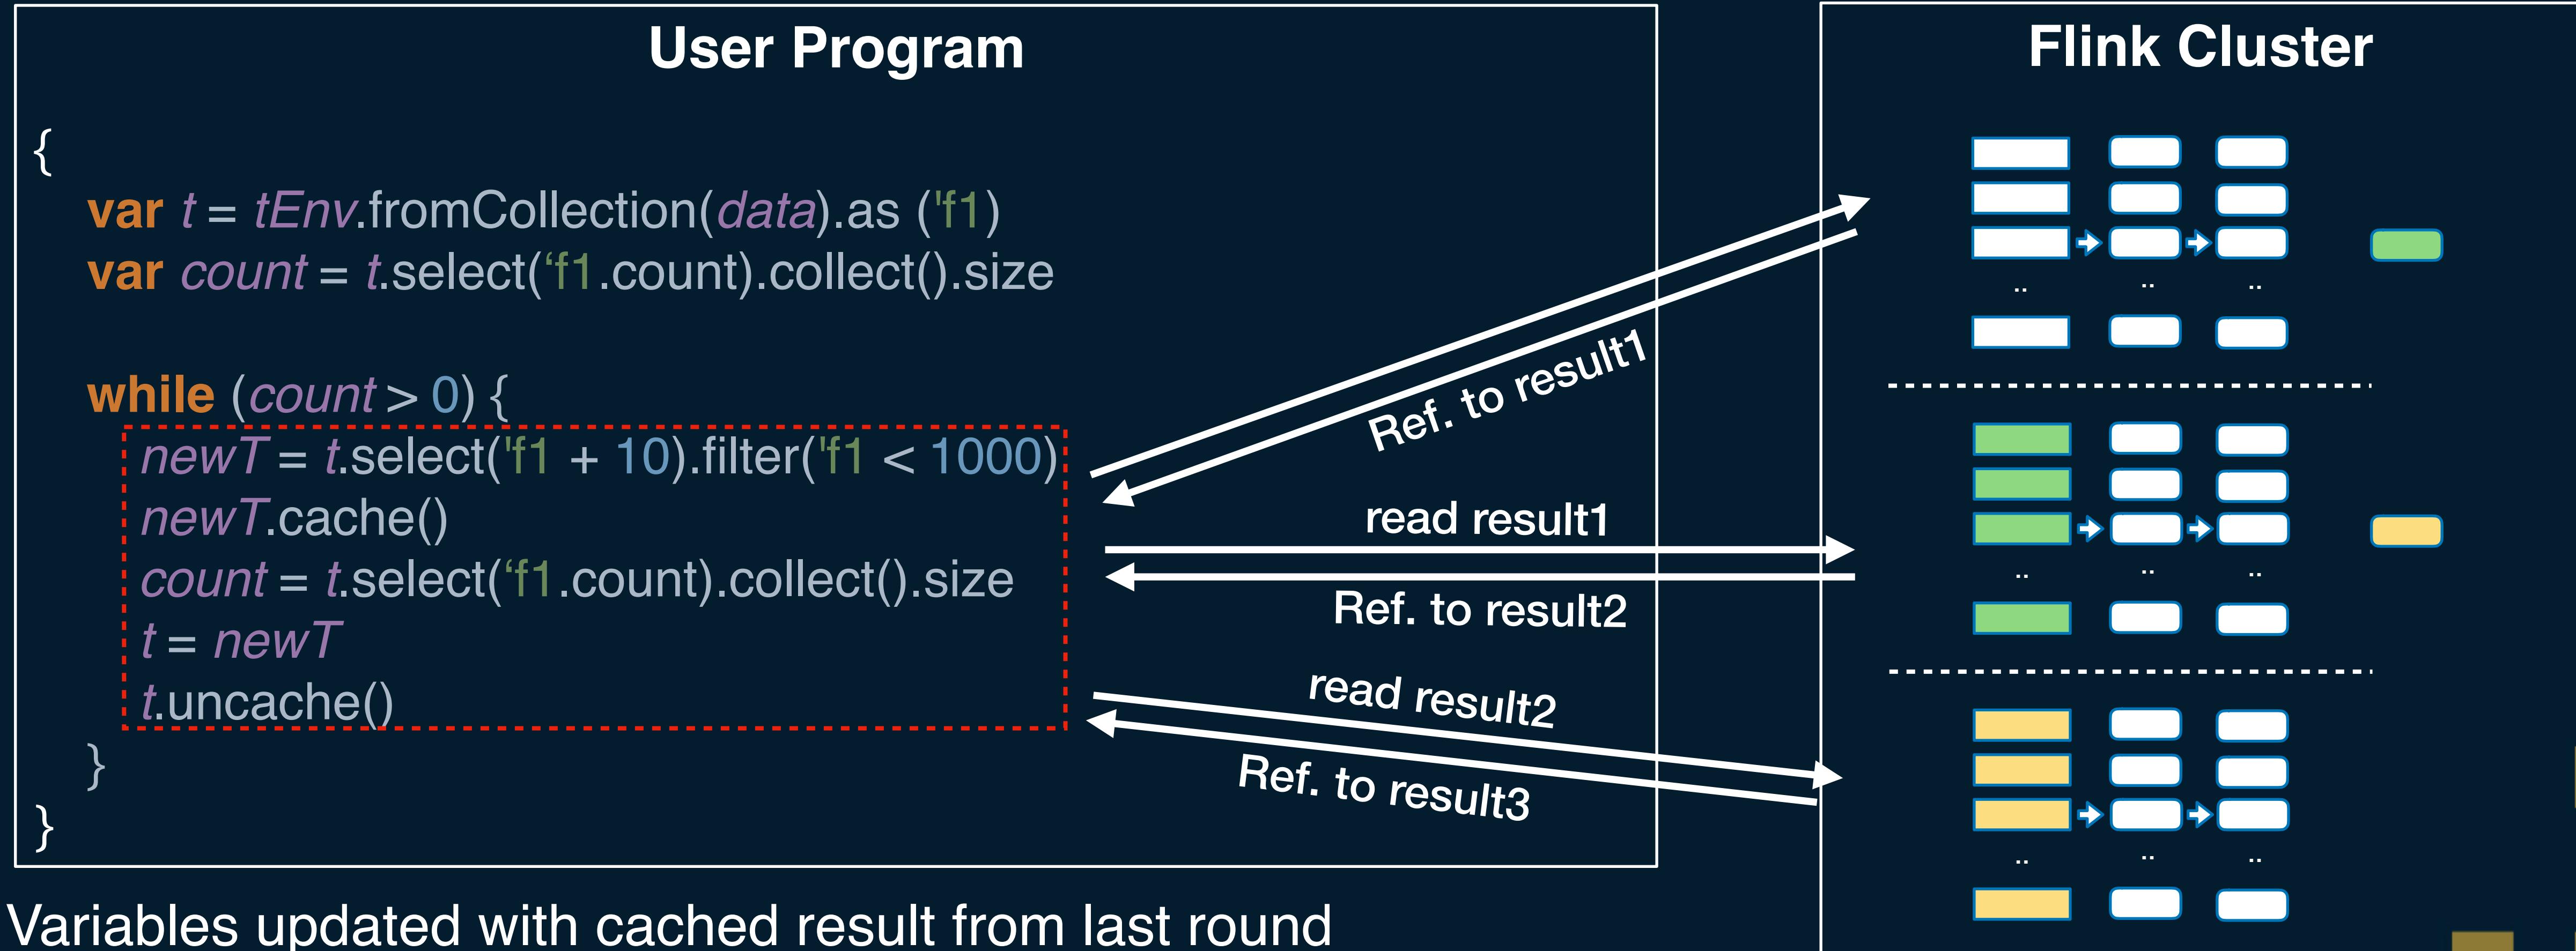
Flink Cluster



Iterative processing with interactive programming



Iterative processing with interactive programming



Iterative processing with interactive programming



The iterations are implemented by the users
迭代由用户实现

Each round of iteration submits a job
每轮迭代提交一个作业

Actively manage the lifecycle of the cached tables
主动cache生命周期维护



Iterative processing with interactive programming



Pros

优点

Simple and intuitive

简单直观

Good flexibility in code logic

代码逻辑灵活性高

Cons

缺点

Overhead of job submission

任务调度开销

Synchronization after each round

轮次间迭代需回到用户程序同步



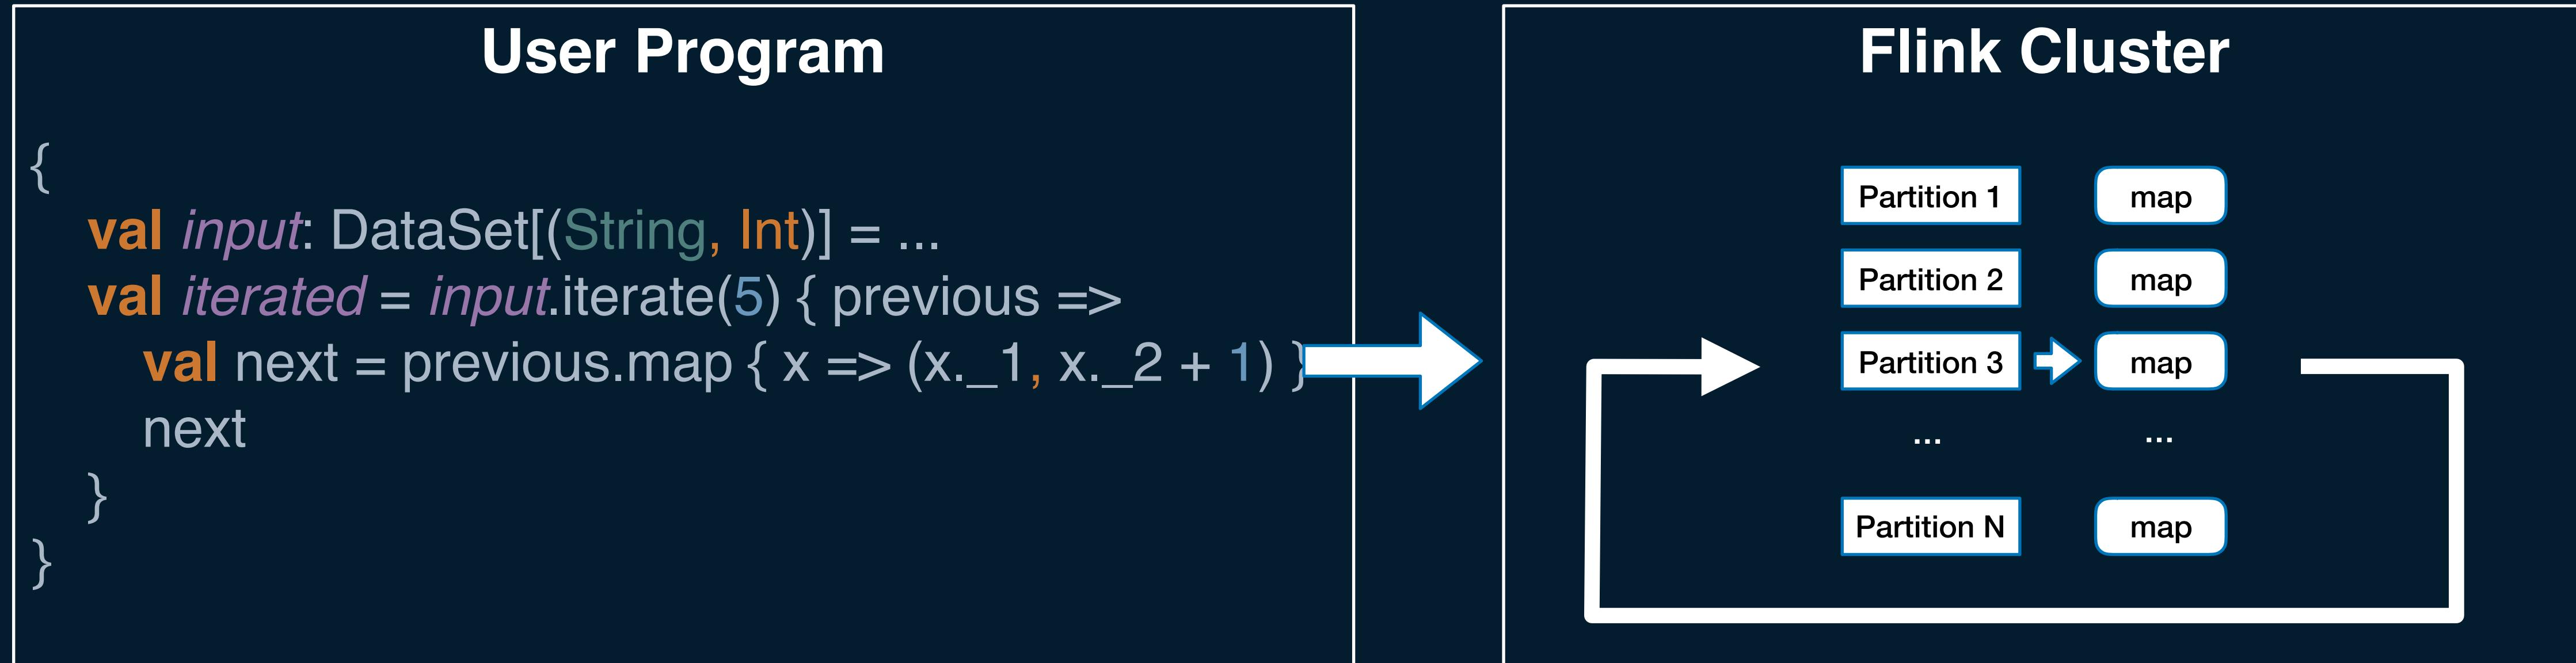
Iteration with descriptive API



- Native iteration implemented by the processing engine
计算引擎实现的原生迭代
- Feedback edge on the processing DAG
在 DAG 基础上增加反馈边
- Exist in Flink DataSet and DataStream
Flink DataSet 和 DataStream 已经采用



Iteration with descriptive API



Single job with a feedback edge
单个有反馈边的作业



Caveats of DataSet and DataStream iteration API



- ▲ No nested iteration support
不支持嵌套迭代
- ▲ No multi-variable iteration support
不支持多变量迭代



The iteration API on Flink Table (Multiple Input)



```
}

val a: Table = ...
val b: Table = ...
val resultSeq = Table.iterate(a, b) {  

    val next_a = b.select('v_b + 1 as 'v_a)  

    val next_b = next_a.select('v_a * 2 as 'v_b)  

    Seq(next_a, next_b)  

}.times(10)  
}
```

Iteration variables

Step function

Termination condition

The iteration API on Flink Table (Nested Iterations)



```
}

val a: Table = ...
val b: Table = ...
val resultSeq = iterate(a, b) {

    val next_a = iterate(a) {

        Seq(a.select('v_a + 1 as 'v_a))
        .times(100).head

        val next_b = next_a.select('v_a * 2 as 'v_b)
        Seq(next_a, next_b)

    }.times(10)
}
```



Custom termination condition

```

}

val a: Table = ...
val b: Table = ...
val termCond: AtomicReference[Table] = ...
val resultSeq = Table.iterate(a, b) {

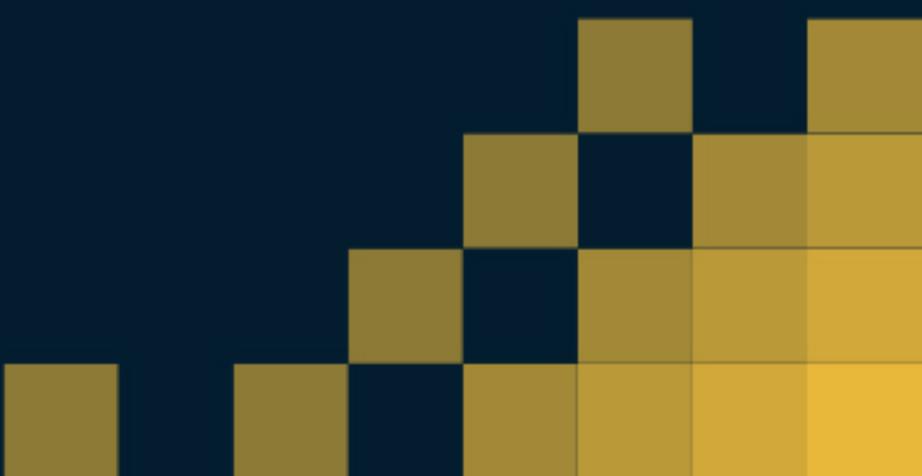
    val next_a = b.select('v_b + 1 as 'v_a)

    val next_b = next_a.select('v_a * 2 as 'v_b)

    [termCond.set(next_a)] ← Stop iteration when next_a is empty
    Seq(next_a, next_b)           当next_a为空时迭代终止

}.untilConverge(termCond.get())
}

```



The iteration API on Flink Table

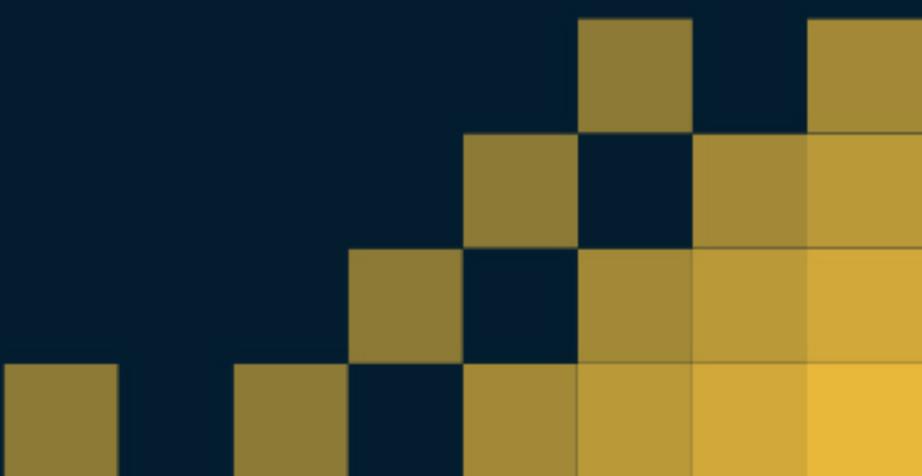


► **About to contribute to the community**
即将贡献至社区

► **Future work**
后续计划

► **Native support for SGD**
增加对 SGD 的原生支持

► **Native support for online learning scenarios**
增加对 online learning 场景的原生支持



AGENDA

摘要



Apache Flink

Brief introduction to Flink Table API

Table API 简介

The API requirements from Machine Learning Algorithms

机器学习算法对 API 的核心需求

The enhancement to Flink Table API

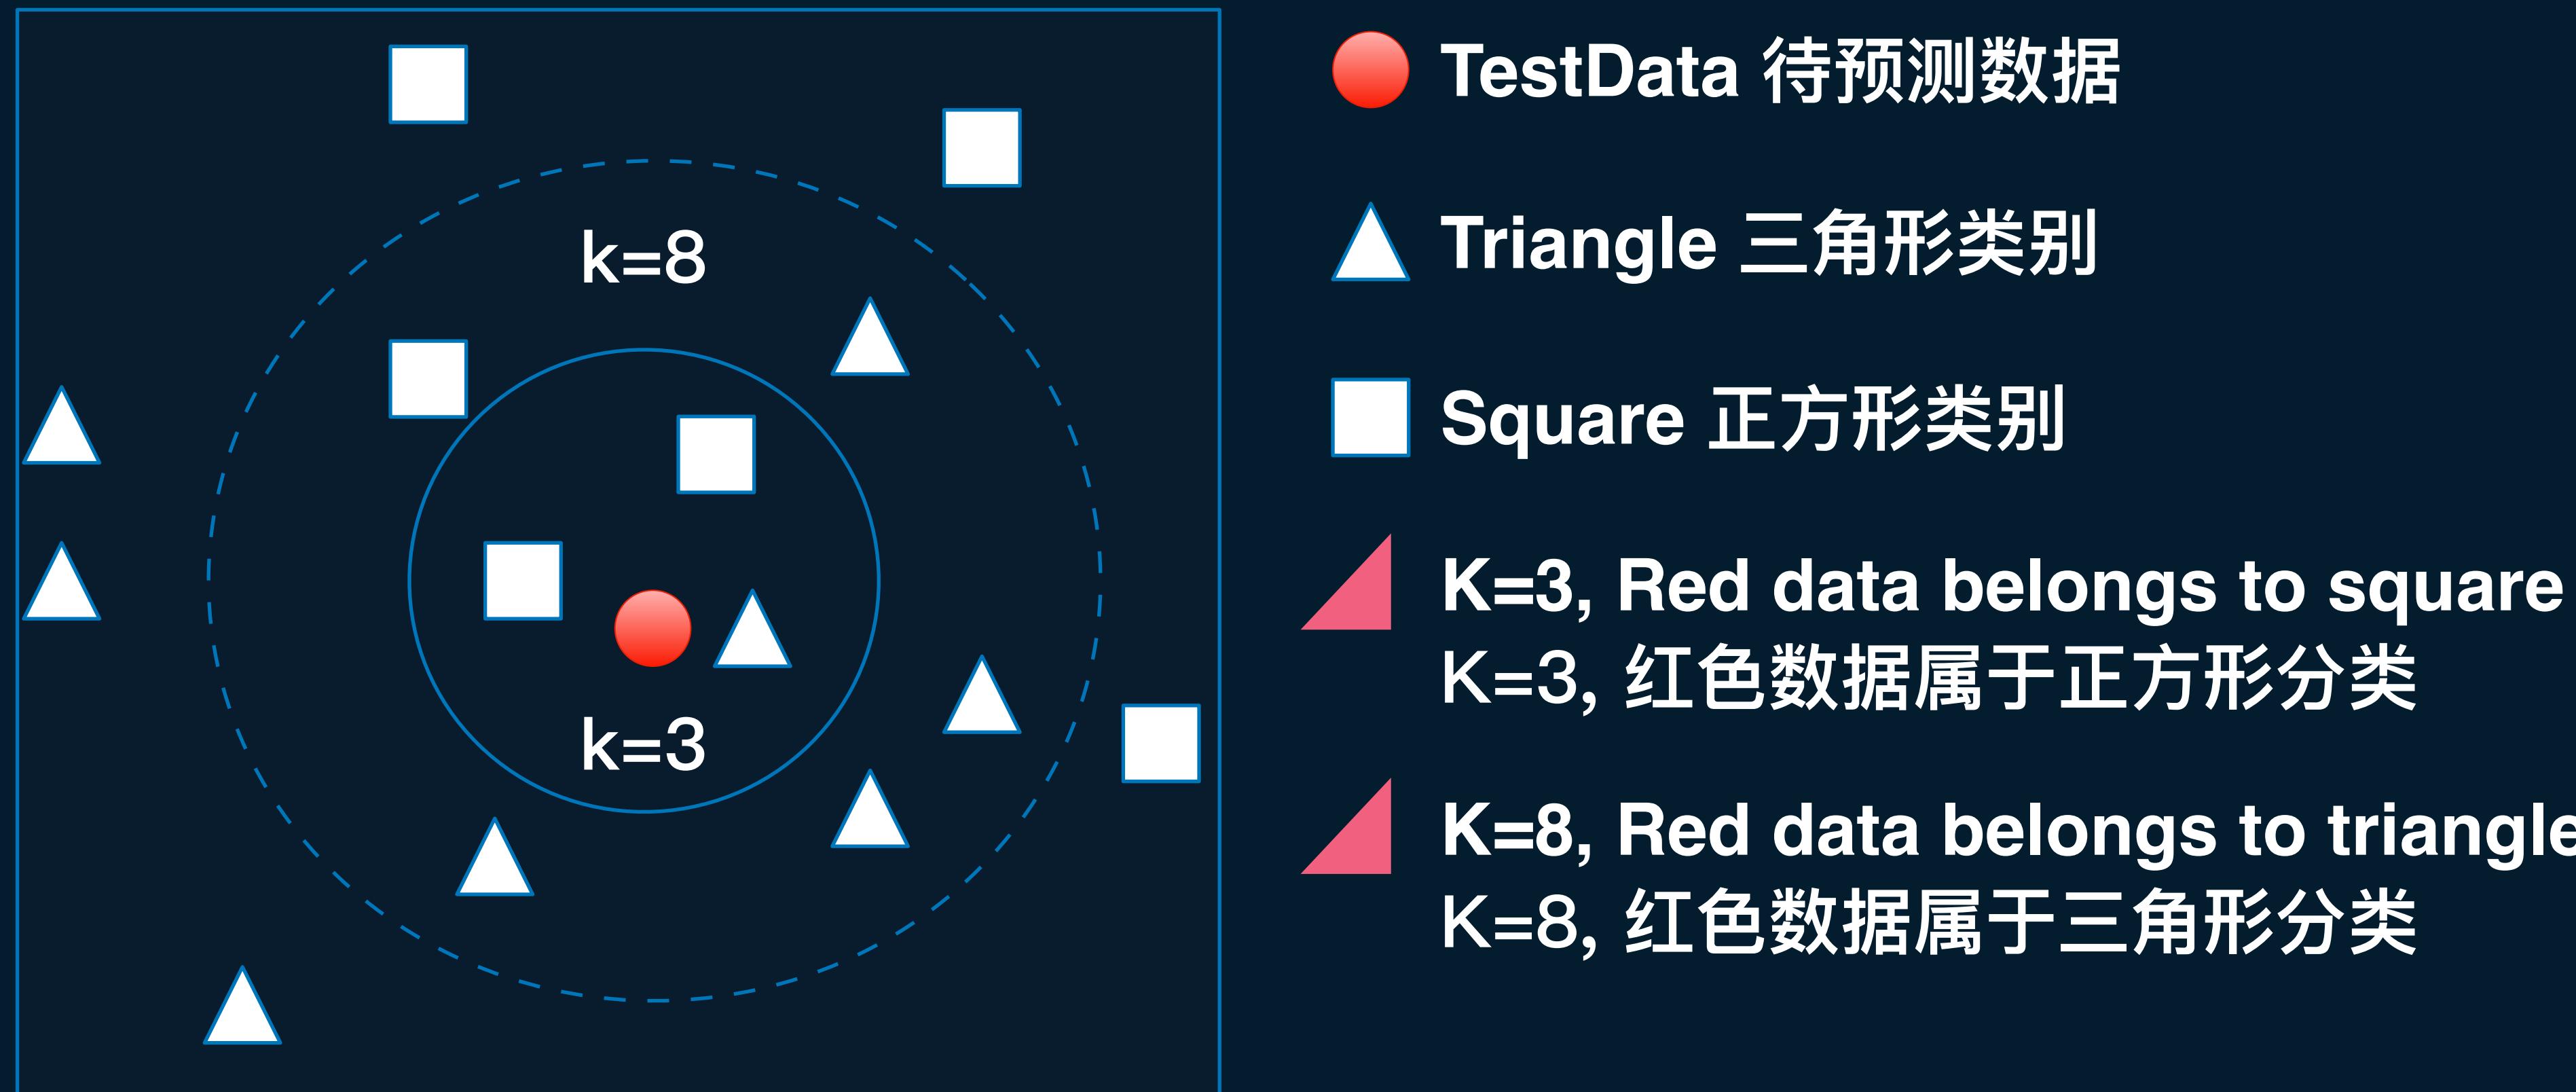
Flink Table API 的扩展

Algorithm examples based on Flink Table API

基于 Flink Table API 的算法实现

KNN(k-nearest neighbor) on TableAPI

Classification of red data in the following figure by KNN algorithm
 KNN 算法对下图红色数据进行分类



KNN(k-nearest neighbor) on TableAPI

// 1. Assign a unique ID to each test data / 每条测试数据分配唯一ID

```
val inputWithId = input.as('t).map(uuidMapFun('t))
```

// 2. Data chunking and concat test and training data/ 数据分块和连接测试和训练数据

```
val inputSplit = TableFlinkMLTools.block(inputWithId, blocks, Some(partitioner))
val crossTuned = trainingSet.as('train).join(inputSplit.as('test), “1=1”)
.select('train, ‘test)
```

// 3. Calculate the distance between the test data and each training data

// 计算测试数与每条训练数据的距离

```
val crossed = crossTuned.flatMap(distanceFlatMapFun('train, ‘test))
.select('singleTrain, 'singleTest, 'id, 'distance)
```

// 4. TopK training data / 计算邻近的K条训练数据

```
val result = crossed.groupBy('id).flatAggregate(topKTableAggFunction(
'singleTrain, 'singleTest, 'id, ‘distance))
```

OutPut: K(3) = (Triangle/三角形, Square/正方形, Square/正方形)



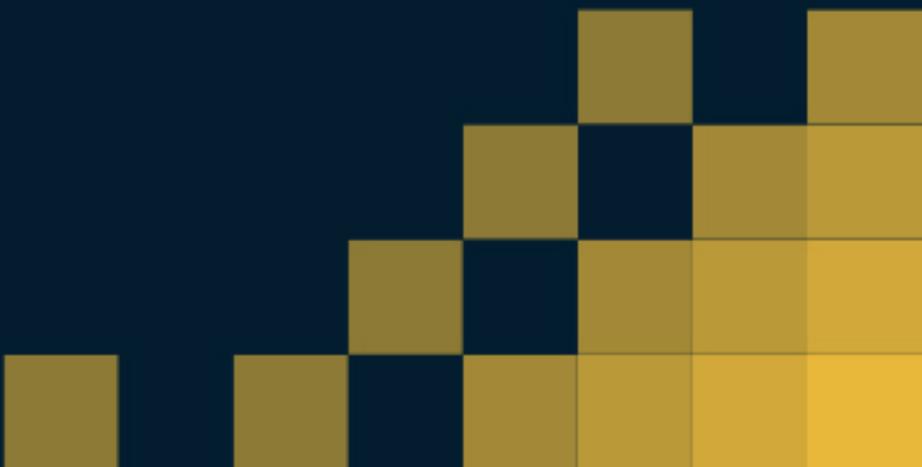
More about Apache Flink

- Source code for KNN on TableAPI
KNN on TableAPI 源代码

Github: <http://t.cn/EU9vPF0>



- Sun Jincheng's Apache Flink album
孙金城 的 Apache Flink 专栏



THANKS

